

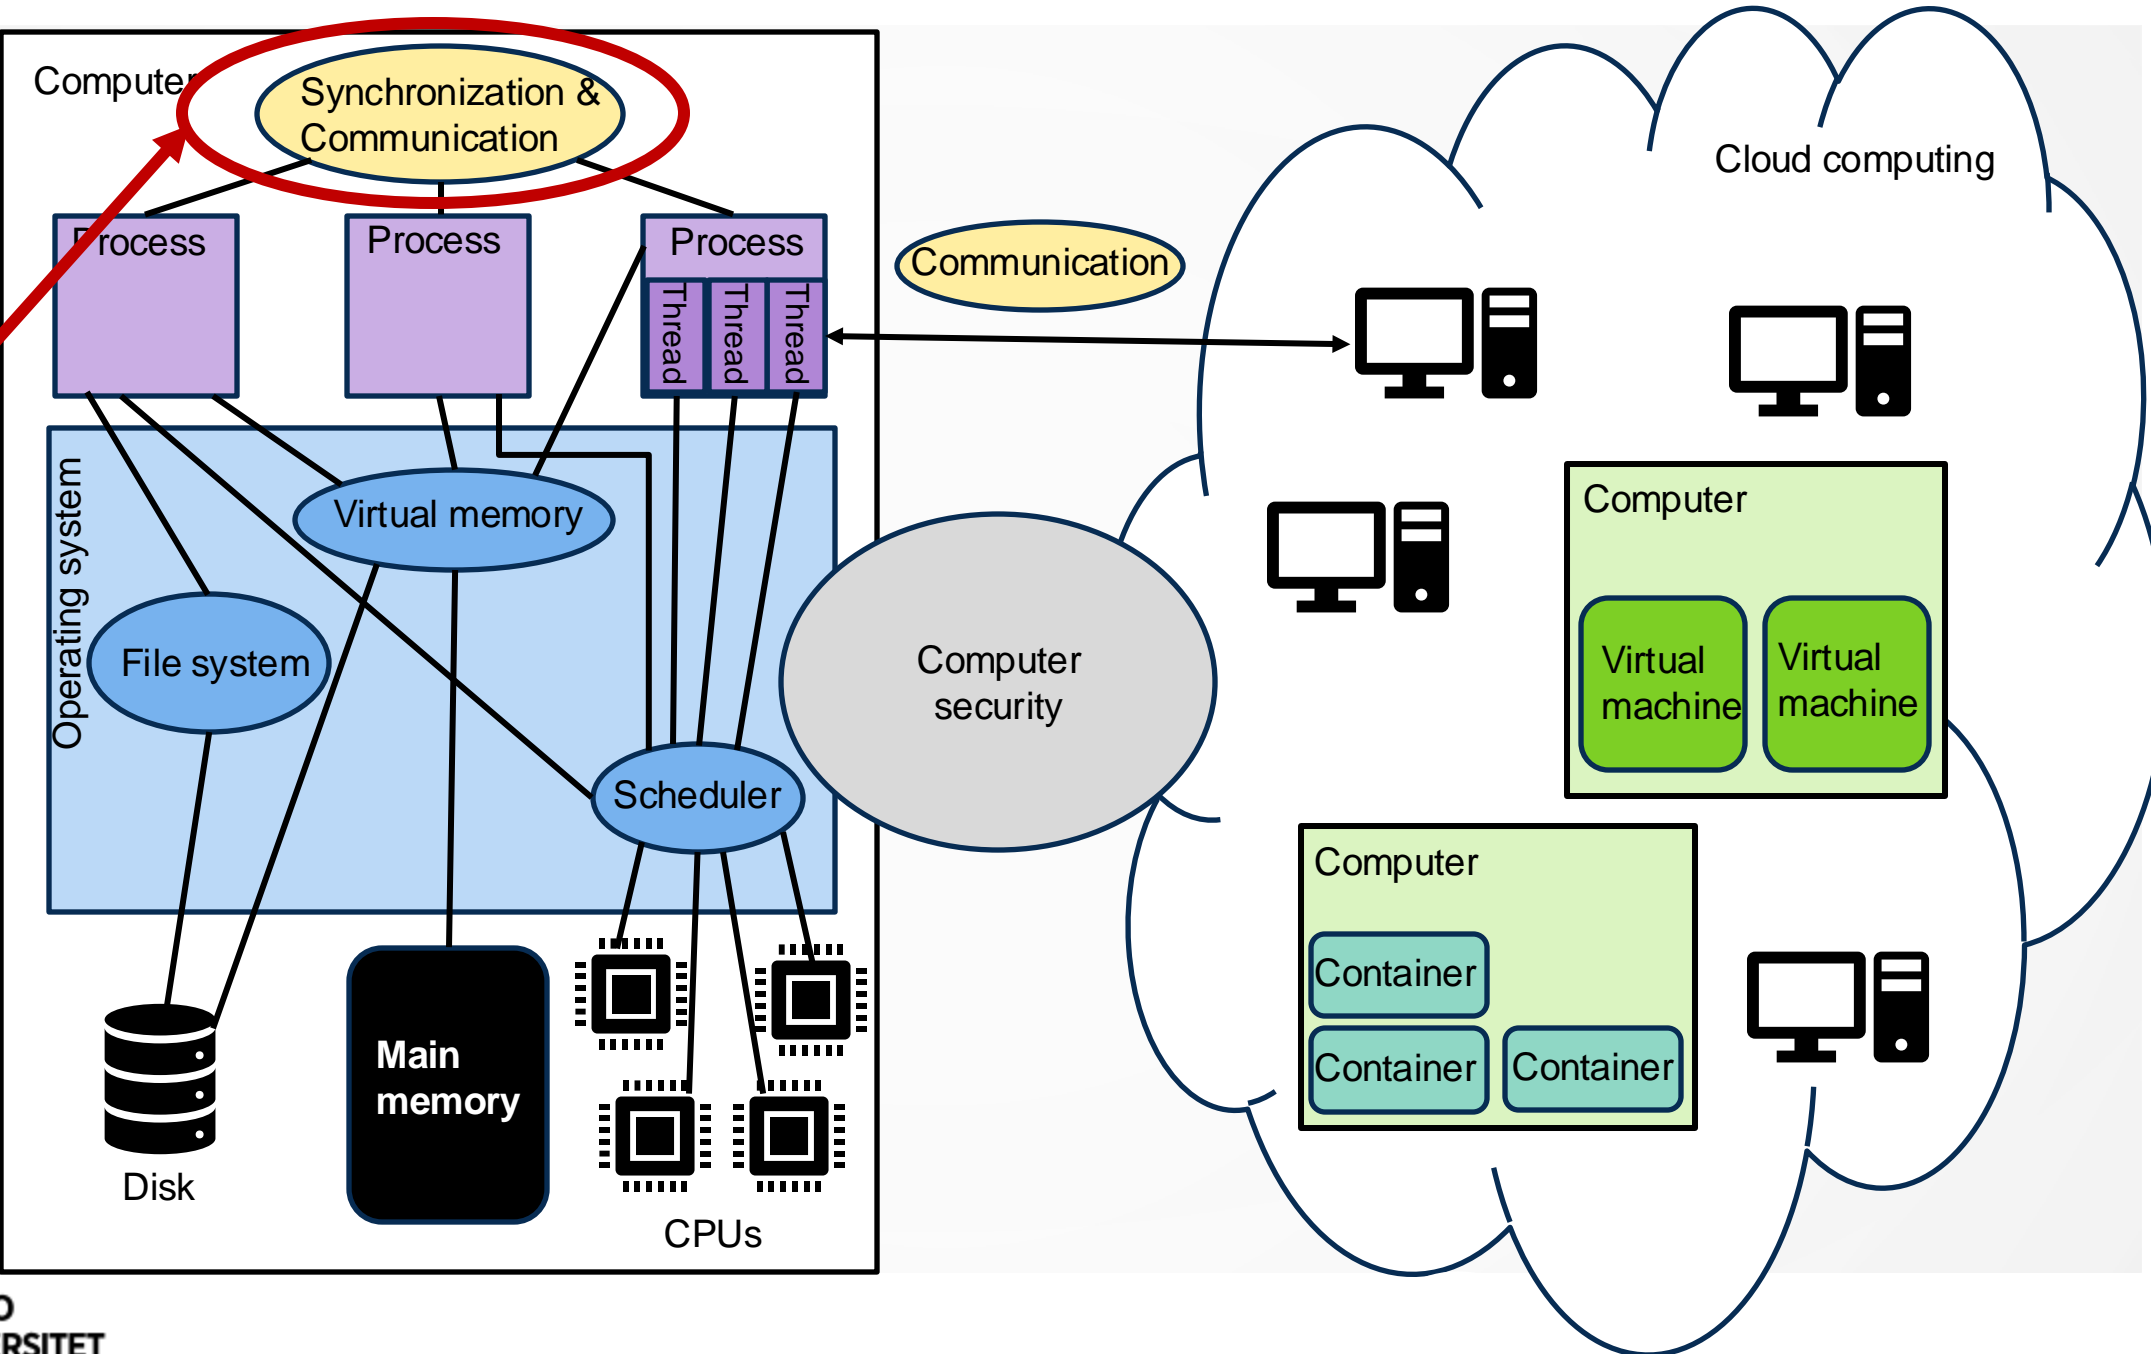


COMPUTING PLATFORMS

Concurrency: Basic concepts and
locks and condition variables



Today's
topic





LEARNING OUTCOMES

- After today's lecture you
 - Know basic concepts related to concurrency such as race conditions and mutual exclusion requirements
 - Understand hardware approaches to supporting mutual exclusion
 - Are able to apply locks and condition variables to simple concurrency problems
 - Can describe and explain producer/consumer problem



EXAMPLE: RACE CONDITION

- Consider the code on the right:
 - X is a variable accessed by two threads (left and right column).
 - The two threads are run concurrently.
 - Single machine instructions can be considered **atomic operations** (no other process can see an intermediate state or interrupt the operation).
 - What are the possible values of X after both threads have executed?

```
shared int X;  
  
load r1=5          load r1=11  
store r1, X        store r1, X
```



EXAMPLE: ANOTHER RACE CONDITION

- Consider the code on the right:
 - X is a variable accessed by two threads (left and right column).
 - The two threads are run concurrently.
 - Single machine instructions can be considered **atomic operations**.
 - What are the possible values of X after both threads have executed?

```
shared int X = 0;

int i = 1;                int j = 1;
while(i <= 2) {           while(j <= 2) {
    X = X + 1;             X = X + 1;
    i = i + 1;             j = j + 1;
}                          }
```



EXAMPLE: ANOTHER RACE CONDITION

X	Thread1	Thread2
0	Load X	
1	Write X+1	
1	Load X	
2	Write X+1	
2		Load X
3		Write X+1
3		Load X
4		Write X+1

X	Thread1	Thread2
0	Load X	
0		Load X
1	Write X+1	
1	Load X	
2	Write X+1	
1		Write X+1
1		Load X
2		Write X+1

X	Thread1	Thread2
0	Load X	
0		Load X
1		Write X+1
1	Write X+1	
1	Load X	
2	Write X+1	
2		Load X
3		Write X+1

...



CONTEXTS OF CONCURRENCY

- Multiple applications
 - **Multiprogramming** invented to allow processing time to be dynamically shared among several active applications
- Structured applications
 - Extension of modular design and structured programming: applications have concurrent processes/threads
- Operating system structure
 - OS implemented as a set of processes/threads



PRINCIPLES OF CONCURRENCY

- **Interleaving** (execution of processes/threads interleaved on a single processor) and **overlapping** (execution of processes/threads happens concurrently on multiple processors)
 - Can be viewed as examples of concurrent processing
 - Both present the same problems – concurrency
- Uniprocessor: relative speed of processes/threads cannot be predicted, depends on
 - Activities of other processes
 - How OS handles interrupts, and
 - Scheduling policies of OS



WHAT IS DIFFICULT TO HANDLE CONCURRENTLY?

- Sharing global resources
- Manage resources optimally
 - No knowledge of who is competing for access
- Difficult to locate programming errors as results are **nondeterministic** and **not reproducible**
 - Errors may only manifest in a rare sequence of events
 - Very hard to reproduce



PROBLEMS COMMONLY OCCURRING IN CONCURRENT PROCESSING

- **Mutual exclusion:**
 - When one process is in a **critical section** accessing shared resources, **no other process can be in a critical section** accessing the same shared resources
- **Deadlock:**
 - Two or more processes are unable to proceed because they are waiting for one of the others to do something
- **Starvation:**
 - A **runnable process is overlooked indefinitely by the scheduler**; although it is able to proceed, it is never chosen



CORRECTNESS OF CONCURRENCY PROBLEM SOLUTIONS

- How do you know that a solution does not work correctly?
 - Find one scenario where the solution does not work (produce a **counter example**)
- How do you know that a solution work correctly?
 - It produces correct result (always, in every scenario)
 - There is no deadlock
 - There is no starvation
 - Prove it: Use formal methods to mathematically prove your point (can be difficult; produces a **proof** of correctness)



CRITICAL SECTIONS AND MUTUAL EXCLUSION

- **Critical section:**
 - Section of code in a process that requires access to shared resources, and that **must not be executed while another process is in a corresponding section of code**
- **Mutex** (=mutual exclusion) solution:
 - Preprotocol: wait is needed, how?
 - Critical section
 - Postprotocol: wake up waiting process

```
/*Process 1*/
```

```
void P1
```

```
{
```

```
    while (true) {
```

```
        /*preceding code*/
```

```
        enter_critical(res
```

```
        /*critical section
```

```
        exit_critical(reso
```

```
        /*following code*/
```

```
    }
```

```
}
```

```
/*Process 2*/
```

```
void P2
```

```
{
```

```
    while (true) {
```

```
        /*preceding code*/
```

```
        enter_critical(resource)
```

```
        /*critical section*/
```

```
        exit_critical(resource)
```

```
        /*following code*/
```

```
    }
```

```
}
```

...



REQUIREMENTS FOR MUTEX SOLUTION

- Mutex must be enforced (functionally correct)
- **Fairness:** Does any thread get a fair chance of entering the critical section?
- **No starvation**
- **No deadlock**
- **Performance:**
 - If there is no contention (i.e. only one thread trying to access the critical section), what is the overhead of acquiring or releasing the mutex?
 - If there are multiple threads on a single CPU trying to access the critical section, are there performance concerns?
 - What about multiple threads on multiple CPUs?



MUTEX: DISABLING INTERRUPTS

- Works only on uniprocessor systems
- Guarantees mutual exclusion
- Waiting happens in the scheduler's Ready-queue
- Disadvantages:
 - Allowing user processes to control interrupts. Can be abused!
 - Efficiency of execution could be noticeably degraded.
 - Does not work in multiprocessor systems.

```
void lock() {  
    disable_interrupts();  
}  
  
void unlock() {  
    enable_interrupts();  
}
```



MUTEX: A FAILED ATTEMPT

- The code on the right does not work correctly. Why?

```
int mutex_flag = 0;

void lock() {
    while (mutex_flag == 1){
    }
    mutex_flag = 1;
}

void unlock() {
    mutex_flag = 0;
}
```



MUTEX: A FAILED ATTEMPT

- The code on the right does not work correctly. Why?
- In `lock()` checking `mutex_flag` and setting it not atomic
- Consider two threads executing `lock()` concurrently:
 - Thread 1 checks that `mutex_flag` is 0
 - Thread 2 checks that `mutex_flag` is 0
 - Thread 1 sets `mutex_flag` to 1
 - Thread 2 sets `mutex_flag` to 1
 - Thus both will enter the critical section at the same time!

```
int mutex_flag = 0;

void lock() {
    while (mutex_flag == 1){
    }
    mutex_flag = 1;
}


void unlock() {
    mutex_flag = 0;
}
```




MUTEX: SPECIAL HARDWARE INSTRUCTIONS

- Compare-and-swap instruction (CAS)
 - Comparison is made between a memory value and a test value
 - If the values are the same, a swap occurs
 - Returns the old memory value
 - Carried out atomically

```
int compare_and_swap(int *word, int testval, int newval) {  
    int oldval = *word;  
    if (oldval == testval) {  
        *word = newval;  
    }  
    return oldval;  
}
```



MUTEX: COMPARE-AND-SWAP

```
int mutex_flag = 0;    /*0=free, 1=locked*/

void lock() {
    while(compare_and_swap(&mutex_flag, 0, 1) == 1) {
        /*do nothing*/
    }
}

void unlock() {
    mutex_flag = 0;
}
```

```
int compare_and_swap(int *word, int testval, int newval) {
    int oldval = *word;
    if (oldval == testval) {
        *word = newval;
    }
    return oldval;
}
```



MUTEX: COMPARE-AND-SWAP

```
int mutex_flag = 0;    /*0=free, 1=locked*/
```

```
void lock() {  
    while(compare_and_swap(&mutex_flag, 0, 1) == 1) {  
        /*do nothing*/  
    }  
}
```

```
void unlock() {  
    mutex_flag = 0;  
}
```

Atomic machine instruction,
returns original value

Locked originally

Set locked (newval)

If open

```
int compare_and_swap(int *word, int testval, int newval) {  
    int oldval = *word;  
    if (oldval == testval) {  
        *word = newval;  
    }  
    return oldval;  
}
```



MUTEX: COMPARE-AND-SWAP

- Advantages
 - Applicable to any number of processes on either uniprocessor or multiprocessor environments
 - Simple and easy to verify
 - Can be used to support multiple critical sections, each with its own variable
- Disadvantages
 - Busy-waiting is employed. Thus a process consumes processor time while waiting
 - Starvation is possible when a process leaves critical section and more than one process is waiting (no FIFO queue)
 - Deadlock is possible



SPINLOCKS

- Compare-and-swap mutex implements a **spinlock**
 - Problem: busy-waiting (process is in a loop waiting for the lock to be released)
 - Spinlocks should only be used for very short critical sections (to minimize busy-waiting time)
- OS can use a spinlock to implement a lock without busy waiting
 - Use a spinlock to protect a critical section that sets the status of the lock (free/locked) and puts threads/processes into Blocked state to wait for the lock to be released
 - When the lock is released, a thread/process can be moved from the Blocked state to the Ready state



LOCKS AND CONDITION VARIABLES

- To implement concurrent programs, we need to solve two problems
 - **Critical section:** At most one thread/process can be inside a critical section at a time
 - **Synchronization:** A thread/process needs to wait for some condition to become true (e.g. data becoming available in a buffer)
- **Locks** solve the critical section problem and **condition variables** solve the synchronization problem
- Many programming languages implement locks and condition variables in a single construct called a **monitor**
 - E.g. Concurrent Pascal, Pascal-Plus, Modula-2, Modula-3, Java, Python (via `threading.Condition` object)
 - Software module consisting of one or more procedures, an initialization sequence, and local data



CRITICAL SECTIONS WITH LOCKS

- Two functions:
 - **Acquire**, lock
 - **Release**, unlock
- A critical section is implemented by acquiring the lock before the critical section and releasing it afterwards

```
lock = threading.Lock()

lock.acquire()
# Critical section
lock.release()
```



SYNCHRONIZATION WITH CONDITION VARIABLES

- Condition variables have no value.
- A condition variable is associated with a FIFO queue of waiting processes
- Condition variables are operated with two functions:
 - **Wait:** suspend the calling process on the given condition
 - **Notify:** resume execution of one process blocked on the given condition (if any)
- Condition variables always associated with a lock to protect the internal state of the condition variable

```
lock = threading.Lock()
condition = threading.Condition(lock)
```

```
# Waiting thread
condition.acquire()
condition.wait()
condition.release()
```

```
# Notifying thread
condition.acquire()
condition.notify()
condition.release()
```




CONDITION VARIABLES: WAIT

- Calling process **always suspends**, process placed in queue
- **Unlocks the lock** associated with the condition variable
 - Allows someone else in the critical section
 - Allows another process awakened from (another?) wait to proceed
- When awakened, the process waits for the lock to proceed
 - May not be the next process to acquire the lock!
 - Condition waited for may not be true anymore, when the process acquires the lock and continues execution – **need to recheck the condition!**



CONDITION VARIABLES: NOTIFY

- Wakes up **first waiting process** if there are any
- Process continues execution, i.e. **it does not release the lock** (Lampson-Redell semantics)
- If no process is waiting in the queue, the notification is lost (no memory)
 - Advanced notifying (with memory) must be handled in some other way (use local variables protected by the lock associated with the condition variable)
- Alternative signaling semantics exist: Signaling process releases the lock and the signaled process continues execution in critical section (Hoare semantics)
 - Most practical implementations use Lampson-Redell semantics



EXAMPLE: PARENT WAITING FOR CHILDREN

- Parent thread creates one or more child threads
- Parent thread wishes to wait for the child thread(s) to finish
 - The parent should sleep until the child threads are done
 - When a child finishes, it should notify the parent thread

```
import threading

def child():
    print("child")
    # notify parent that I am done
```

```
def parent():
    print("parent start")
    child1 = threading.Thread(target=child)
    child2 = threading.Thread(target=child)
    child1.start()
    child2.start()
    # wait for children to be done
    print("parent end")
```



EXAMPLE: PARENT WAITING FOR CHILDREN

```
import threading

count = 0
lock = threading.Lock()
done = threading.Condition(lock)

def child():
    global count
    print("child")
    done.acquire()
    count -= 1
    if count == 0:
        done.notify()
    done.release()
```

```
def parent():
    global count
    count = 2
    print("parent start")
    child1 = threading.Thread(target=child)
    child2 = threading.Thread(target=child)
    child1.start()
    child2.start()
    done.acquire()
    while count != 0:
        done.wait()
    done.release()
    print("parent end")
```



PRODUCER/CONSUMER PROBLEM

- General situation:
 - One (or more) producer is generating data items and placing them in a buffer (one at a time)
 - One (or more) consumer is taking items out of the buffer (one at a time)
 - Producers and consumers may access the buffer concurrently
- Problem: How to ensure that
 - ... producer must wait for consumer if buffer is full, and
 - ... consumer must wait for producer if buffer is empty
- Who communicates with whom? How? (Communication problem)
- Who waits for whom? When? How? (Synchronization problem)



PRODUCER/CONSUMER: INFINITE BUFFER

- First, assume infinite buffer (cannot overflow!)
 - Producer does not need to check if there are free slots in the buffer
- We will implement this with a Python list where we always can append a new element



PRODUCER/CONSUMER PROBLEM: INFINITE BUFFER

```
import threading

buffer = []
lock = threading.Lock()
notempty = threading.Condition(lock)

def producer():
    global buffer
    while(True):
        item = produce_item()
        lock.acquire()
        buffer.append(item)
        notempty.notify()
        lock.release()
```

```
def consumer():
    global buffer
    while(True):
        lock.acquire()
        while len(buffer) == 0:
            notempty.wait()
        item = buffer.pop(0)
        lock.release()
        consume_item(item)
```



PRODUCER/CONSUMER PROBLEM: INFINITE BUFFER

```
import threading

buffer = []
lock = threading.Lock()
notempty = threading.Condition(lock)

def producer():
    global buffer
    while(True):
        item = produce_item()
        lock.acquire()
        buffer.append(item)
        notempty.notify()
        lock.release()
```

```
def consumer():
    global buffer
    while(True):
        lock.acquire()
        while len(buffer) == 0:
            notempty.wait()
        item = buffer.pop(0)
        lock.release()
        consume_item(item)
```

Protect critical section

Notify: there are now items in the buffer

Wait until buffer not empty
Why use while, not if?



PRODUCER/CONSUMER: FINITE BUFFER

- Block:
 - Insert in full buffer (producer waits)
 - Remove from empty buffer (consumer waits)
- Unblock
 - Item inserted (producer notifies consumer)
 - Item removed (consumer notifies producer)



PRODUCER/CONSUMER PROBLEM: FINITE BUFFER

```
import threading

buffer = []
lock = threading.Lock()
notempty = threading.Condition(lock)
notfull = threading.Condition(lock)

def producer():
    global buffer
    while(True):
        item = produce_item()
        lock.acquire()
        while len(buffer) == 10:
            notfull.wait()
        buffer.append(item)
        notempty.notify()
        lock.release()
```

```
def consumer():
    global buffer
    while(True):
        lock.acquire()
        while len(buffer) == 0:
            notempty.wait()
        item = buffer.pop(0)
        notfull.notify()
        lock.release()
        consume_item(item)
```



PRODUCER/CONSUMER PROBLEM: FINITE BUFFER

```
import threading

buffer = []
lock = threading.Lock()
notempty = threading.Condition(lock)
notfull = threading.Condition(lock)

def producer():
    global buffer
    while(True):
        item = produce_item()
        lock.acquire()
        while len(buffer) == 10:
            notfull.wait()
        buffer.append(item)
        notempty.notify()
        lock.release()
```

```
def consumer():
    global buffer
    while(True):
        lock.acquire()
        while len(buffer) == 0:
            notempty.wait()
        item = buffer.pop(0)
        notfull.notify()
        lock.release()
        consume_item(item)
```

Wait until buffer not full

Why use while, not if?

Notify: the buffer is not full

Wait until buffer not empty

Why use while, not if?

Notify: there are now items in the buffer



SUMMARY

- Critical sections and mutual exclusion (mutex)
- Spinlocks with hardware assistance
- Locks and condition variables
 - Locks implement mutual exclusion
 - Condition variables implement synchronization
- Producer/consumer problem