

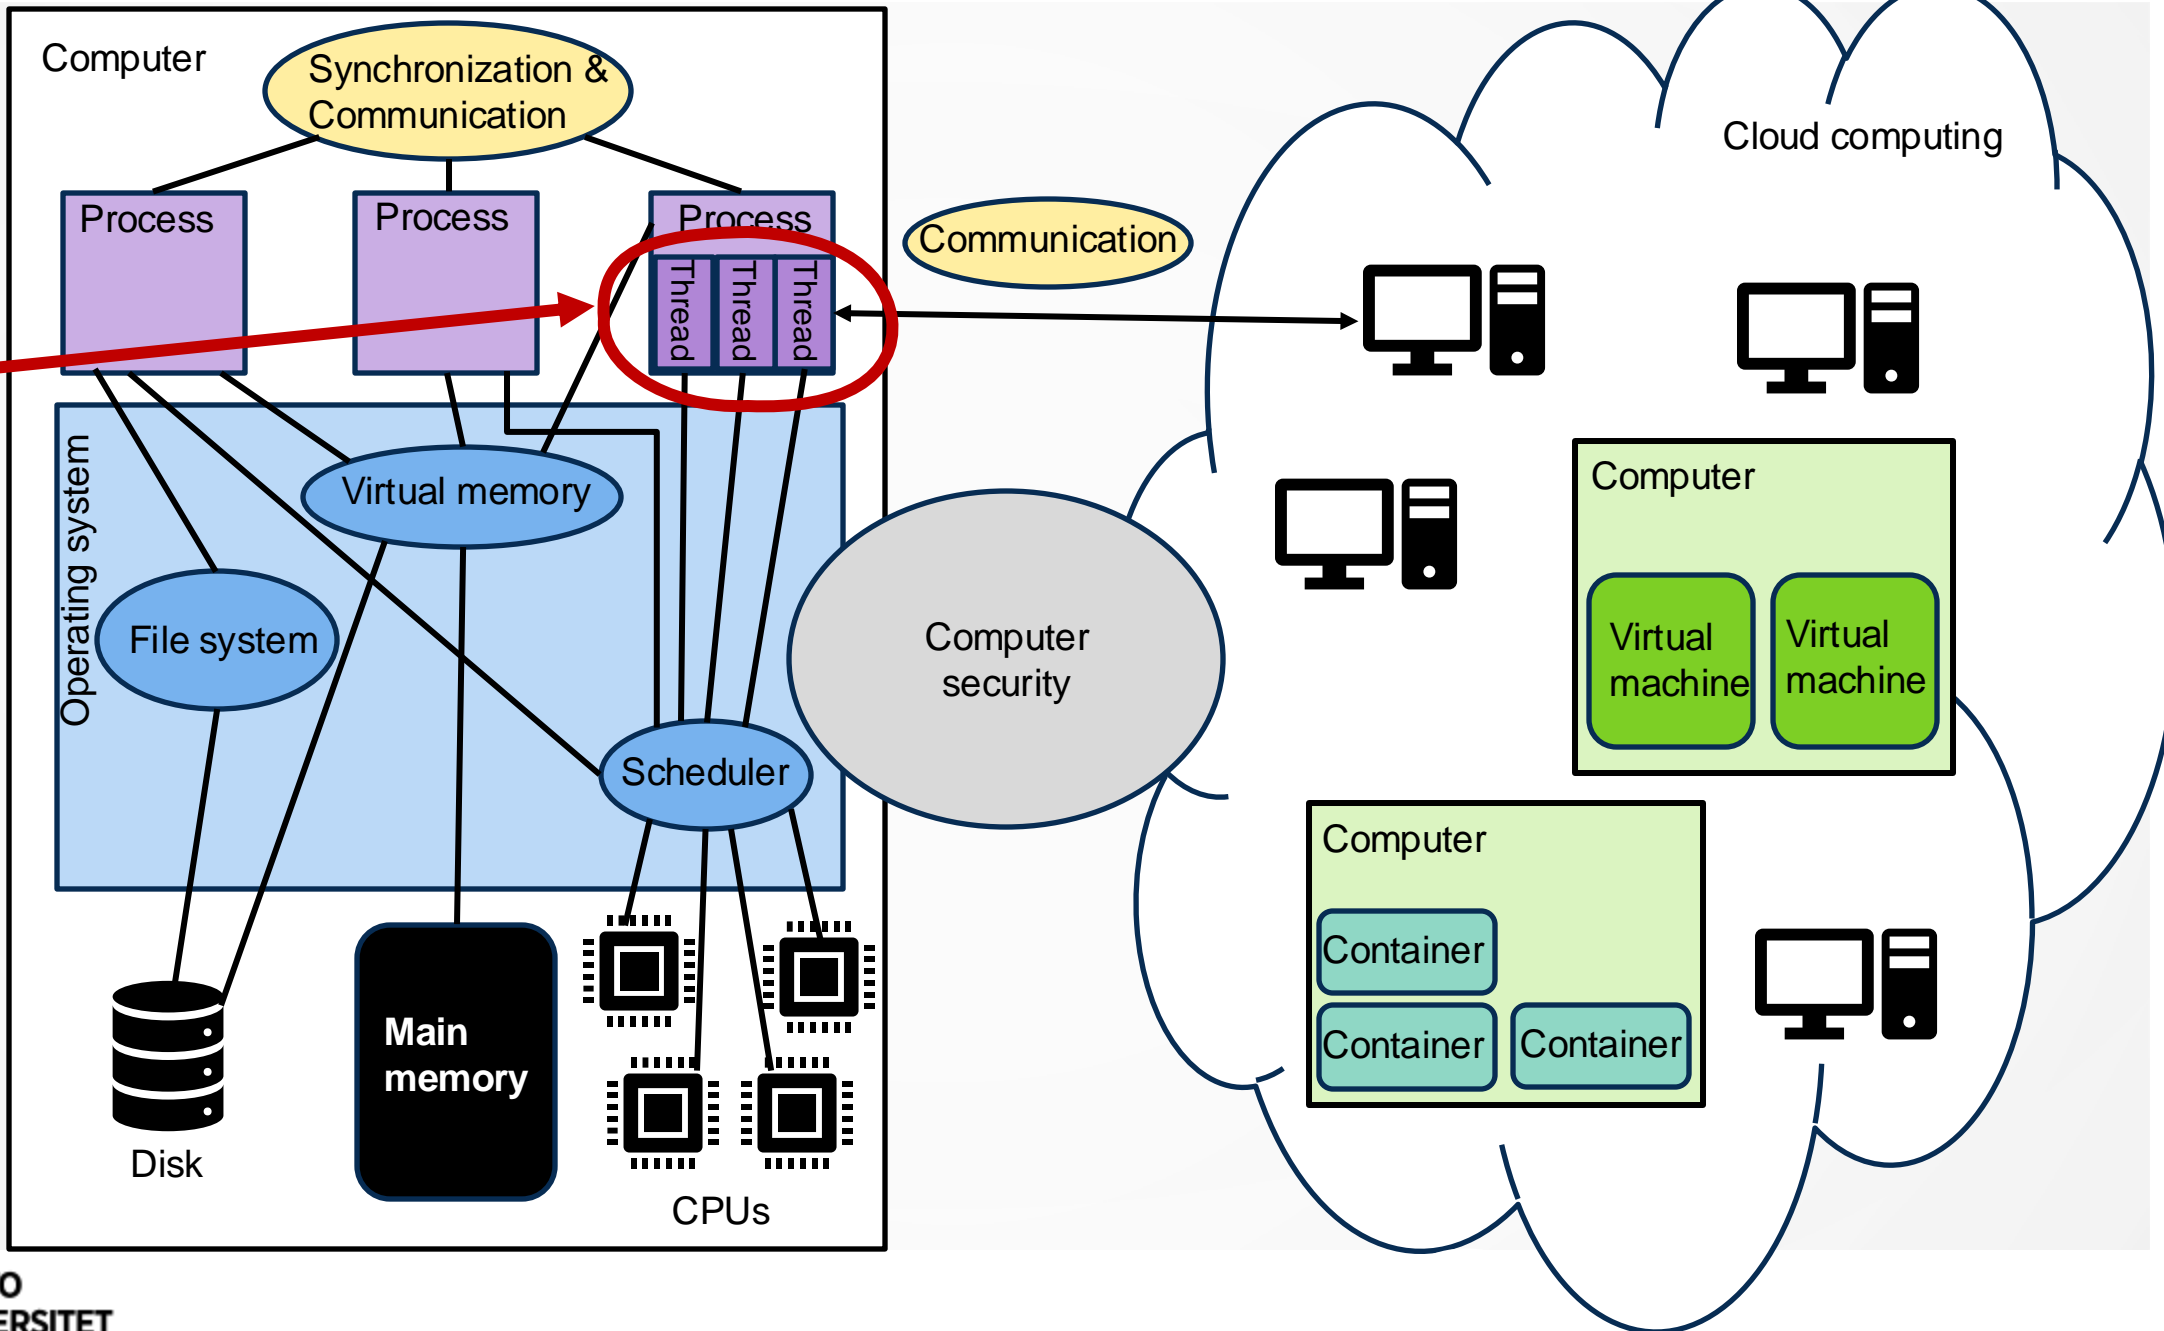


COMPUTING PLATFORMS

Threads



Today's
topic





LEARNING OUTCOMES

- After today's lecture, you
 - Understand the relationship and differences between threads and processes
 - Are able to describe how user-level and kernel-level threads differ and provide a comparison of the advantages and disadvantages
 - Know of Amdahl's law and how that relates to multithreading



WHAT ARE THREADS: TWO PROCESS CHARACTERISTICS

- Resource ownership
 - Own virtual address space holding the process image
 - OS protects processes from each other to prevent unwanted interference wrt. resources
- Scheduling / execution
 - Process has an execution state (**running**, **ready**, etc.) and a dispatching priority
 - Entity that is scheduled and dispatched by the OS



WHAT ARE THREADS: TWO PROCESS CHARACTERISTICS

Process / task

- Resource ownership
 - Own virtual address space holding the process image
 - OS protects processes from each other to prevent unwanted interference wrt. resources
- Scheduling / execution

Thread / lightweight process

- Process has an execution state (**running**, **ready**, etc.) and a dispatching priority
 - Entity that is scheduled and dispatched by the OS
- **Multithreading:** Ability of an operating system to support **multiple, concurrent** paths of execution **within a single process**



WHY THREADS?

- What are advantages of threads over processes?
- What are the uses of threads?



ADVANTAGES OF THREADS

- **Faster to create** a new thread than a new process: Can be ten times faster
- **Faster to terminate** a thread than a process
- **Switching** between two threads **faster** than switching between processes
- Threads **enhance efficiency in communication between programs** (when implemented as threads, not separate processes)
 - Communication between processes in most cases requires the intervention of the kernel (OS protects processes from each other)
 - Threads within a process share memory and files, and thus can communicate with each other **without invoking the kernel**



USES OF THREADS

- **Foreground and background work**
 - E.g., a spreadsheet program: one thread displays user interface, another thread executes user commands and updates the spreadsheet
- **Asynchronous processing**
 - E.g., a word processor may create a backup file of unsaved changes once every minute using a dedicated thread
- **Speed of execution**
 - One thread can be reading data from disk, while another is computing on a previously read batch of data.
 - On a multiprocessor system, multiple threads of the same process can execute simultaneously
- **Modular program structure**



PROCESS MODEL WITH THREADS

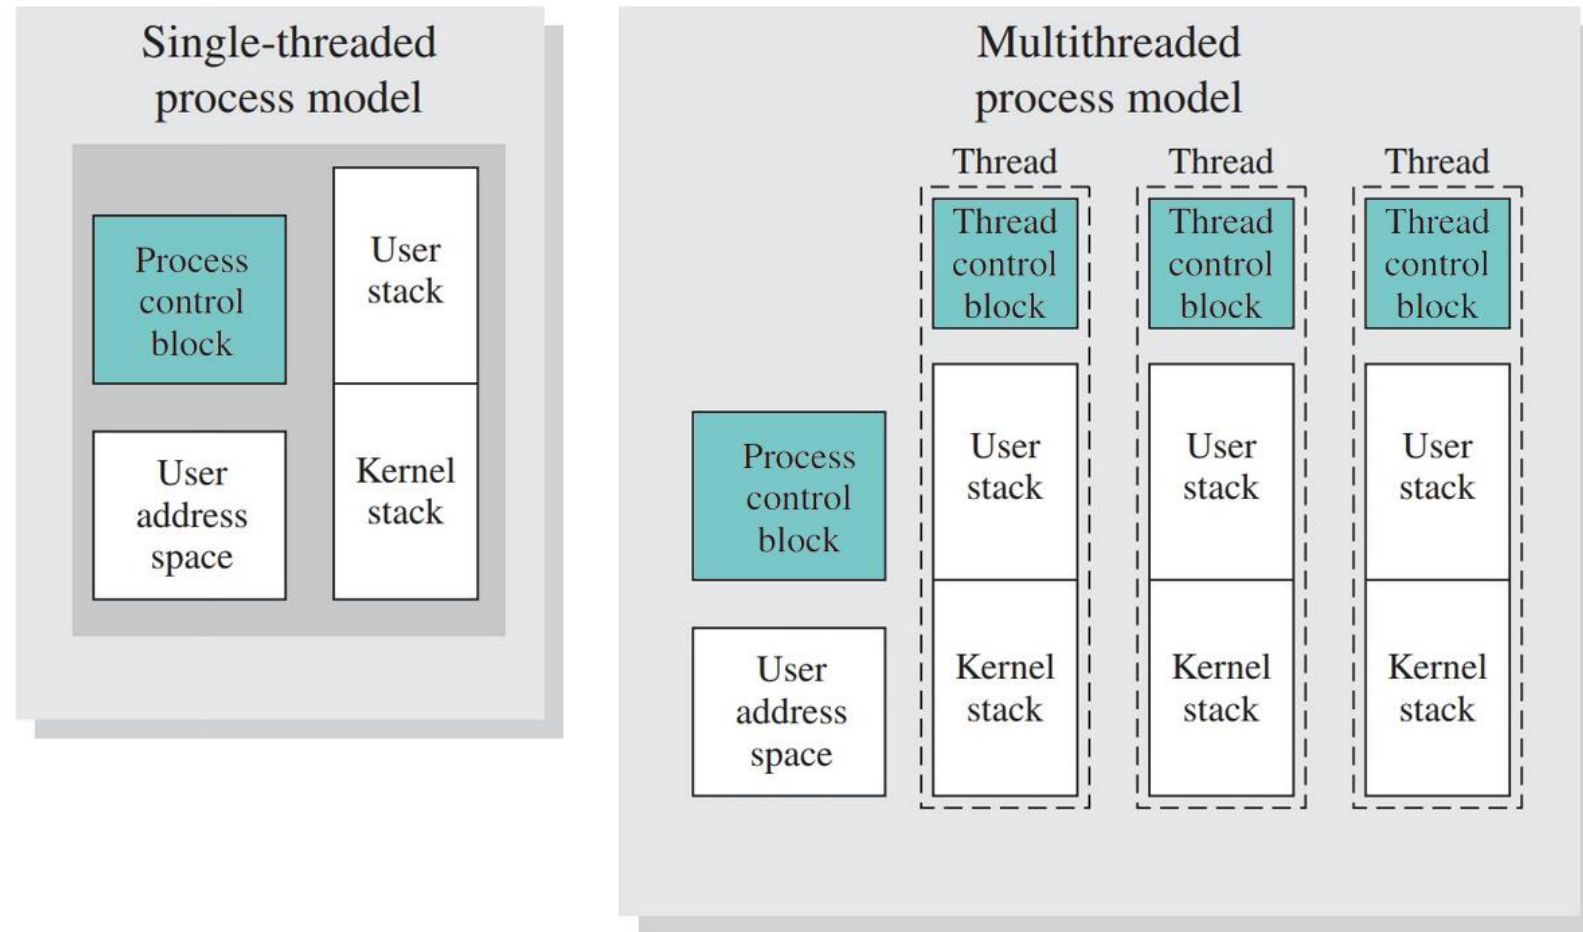


Figure 4.2 Single-Threaded and Multithreaded Process Models

Figure from [Stallings, Operating systems: Internals and design principles, 9th ed]



EACH THREAD HAS...

- **Private:**
 - Execution state (**running**, **ready**, ...)
 - Saved thread context when not running (CPU registers)
 - Execution stack
 - Some static storage for local variables, i.e., "global" data for a thread
- **Shared:**
 - Access to the **memory** and **resources** of its process
 - When one thread alters an item of data in memory, all threads in the process see the results (if they access it).
 - If one thread opens a file with read privileges, all threads (in the same process) can also read from that file. (*What about writing...*)

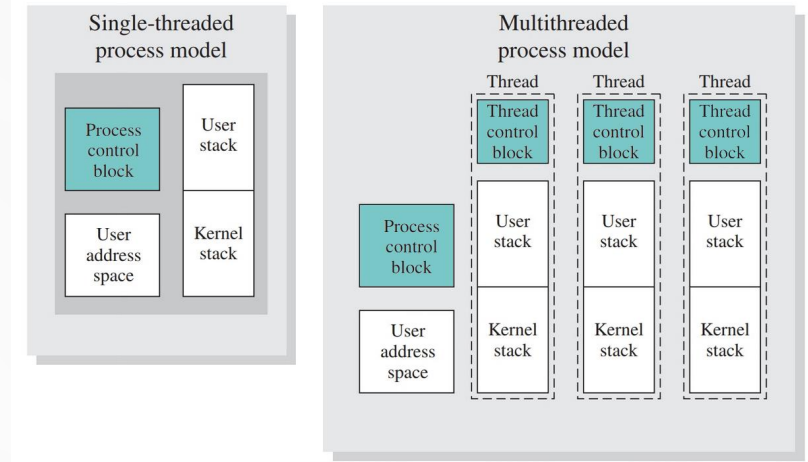



Figure 4.2 Single-Threaded and Multithreaded Process Models



THREAD STATES

- **Scheduling** and **dispatching** is done **on a thread basis**
- Most of the state information dealing with execution is maintained in thread-level data structures
 - **Suspending** a process suspends all its threads
 - **Terminating** a process terminates all its threads
- Thread states are thus
 - **Running**
 - **Ready**
 - **Blocked**





EXAMPLE WITH THREADS

- Thread C begins to run after thread A exhausts its time quantum. Note that thread B is also ready to run. This is a **scheduling decision** (a topic covered in the next lecture)

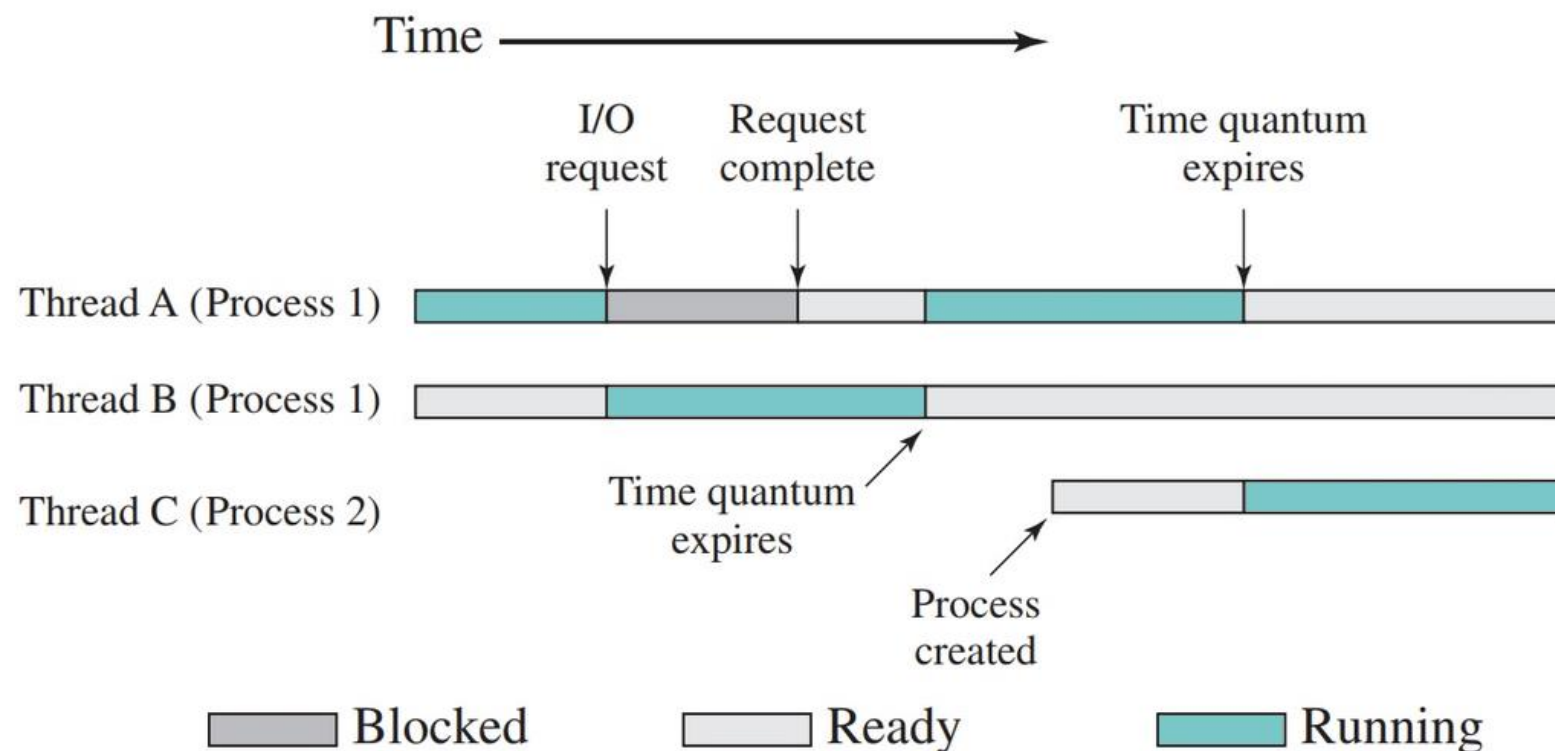


Figure 4.4 Multithreading Example on a Uniprocessor



THREAD SYNCHRONIZATION

- All threads of a process **share the same address space** and other **resources** (e.g. open files)
- Any alteration of data in memory or other resources **affects all threads**
- Necessary to **synchronize** the threads so that they do not interfere with each other or corrupt data structures
- Synchronization issues covered later in the course



USER-LEVEL THREADS VS KERNEL-LEVEL THREADS

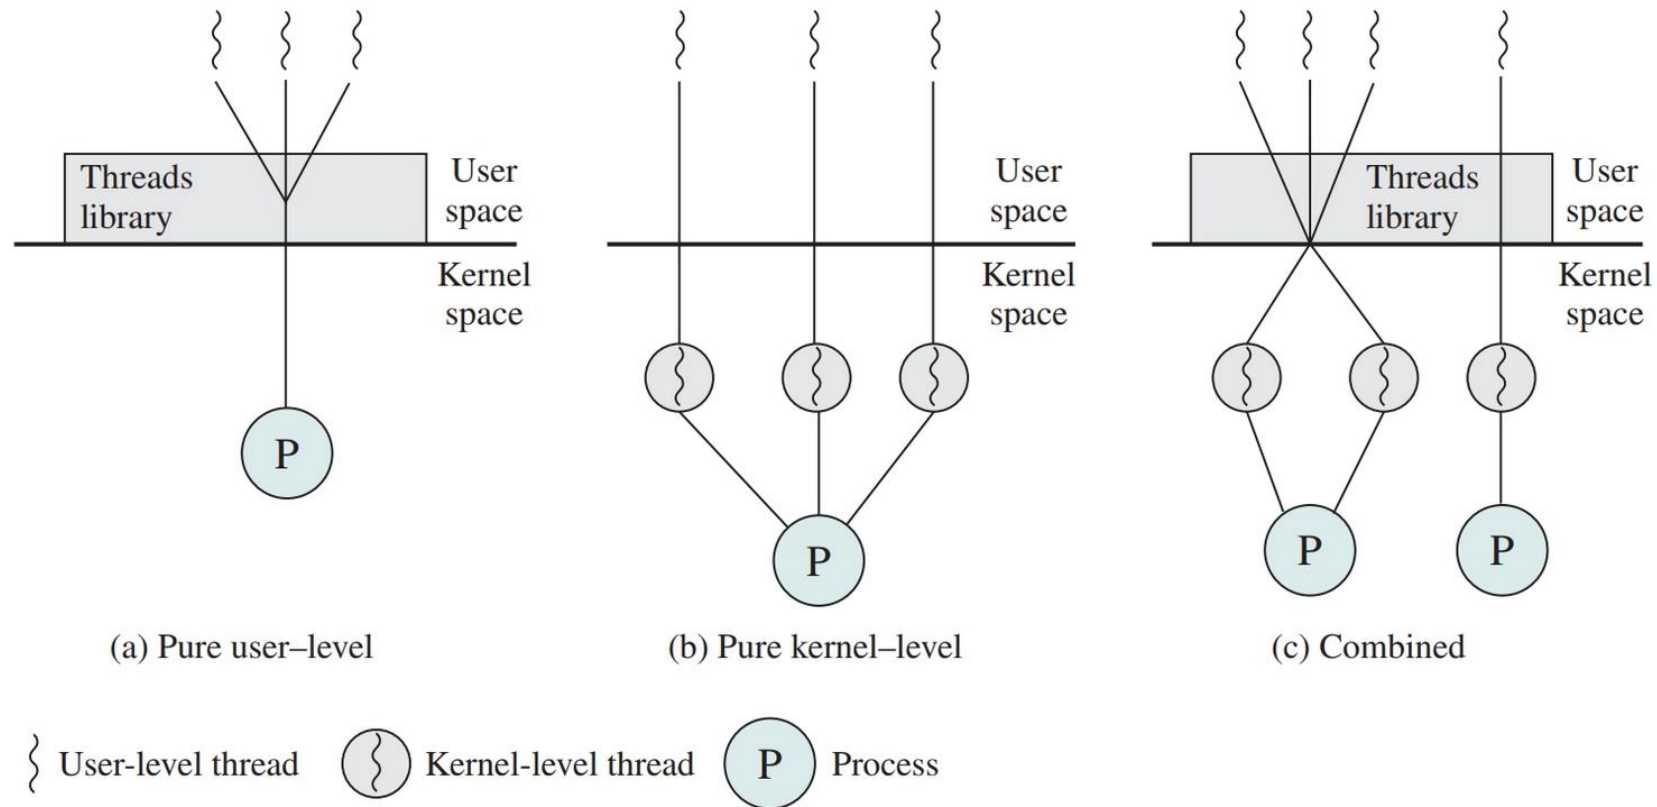



Figure 4.5 User-Level and Kernel-Level Threads



USER-LEVEL THREADS (ULT)

- Thread management is done by the application;
The **kernel is not aware of the existence of threads**.
- Multithreaded applications are programmed using a **threads library** which contains code for
 - Creating and destroying threads
 - Passing messages and data between threads
 - Scheduling thread execution
 - Saving and restoring thread contexts
- Kernel **schedules the process as a unit** and assigns a single execution state to the process (**Ready, Running, Blocked,...**)






ULT – ADVANTAGES AND DISADVANTAGES

- Thread switching does not require kernel-mode privileges: saves the overhead of two mode switches
- Scheduling can be application specific
- ULTs can run on any OS: The kernel does not even know of the existence of threads
- A **blocking system call** in ULT causes all threads in the process to be blocked
- A multithreaded application cannot take advantage of multiprocessing: Kernel assigns one process to one processor at a time



KERNEL-LEVEL THREADS (KLT)

- All thread management done by the kernel;
No thread management code in the application level
- Kernel maintains context information for the process as a whole and for each individual thread in the process
- Scheduling in the OS done on a thread basis
- Overcomes two weaknesses of ULT:
 - Kernel can schedule multiple threads from the same process on multiple processors
 - If one thread is blocked, other threads in the process can still be scheduled by the OS
- Kernel can be multithreaded too!
- Drawback: Thread switching now requires two mode switches; much slower than for ULT





MULTITHREADING ON MULTICORE SYSTEMS

- In a multicore system many threads of a single application can run concurrently
- How does this affect performance?
 - Depends how well the application can exploit the parallel resources



AMDAHL'S LAW

$$\text{speedup} = \frac{\text{time to execute program on a single processor}}{\text{time to execute program on } N \text{ parallel processors}}$$


- Let f be the fraction of code that is (infinitely) parallelizable
- $(1-f)$ is then the fraction of code that is inherently serial

$$\text{speedup} = \frac{1}{(1 - f) + \frac{f}{N}}$$




AMDAHL'S LAW

$$\text{speedup} = \frac{1}{(1 - f) + \frac{f}{N}}$$



(a) Speedup with 0%, 2%, 5%, and 10% sequential portions



(b) Speedup with overheads

Figure 4.7 Performance Effect of Multiple Cores



SUMMARY

- **Thread** is a scheduling unit
- Two ways of implementing threads: **user-level threads** and **kernel-level threads**
- Benefits of multithreading depend on the fractions of serial and parallelizable code