

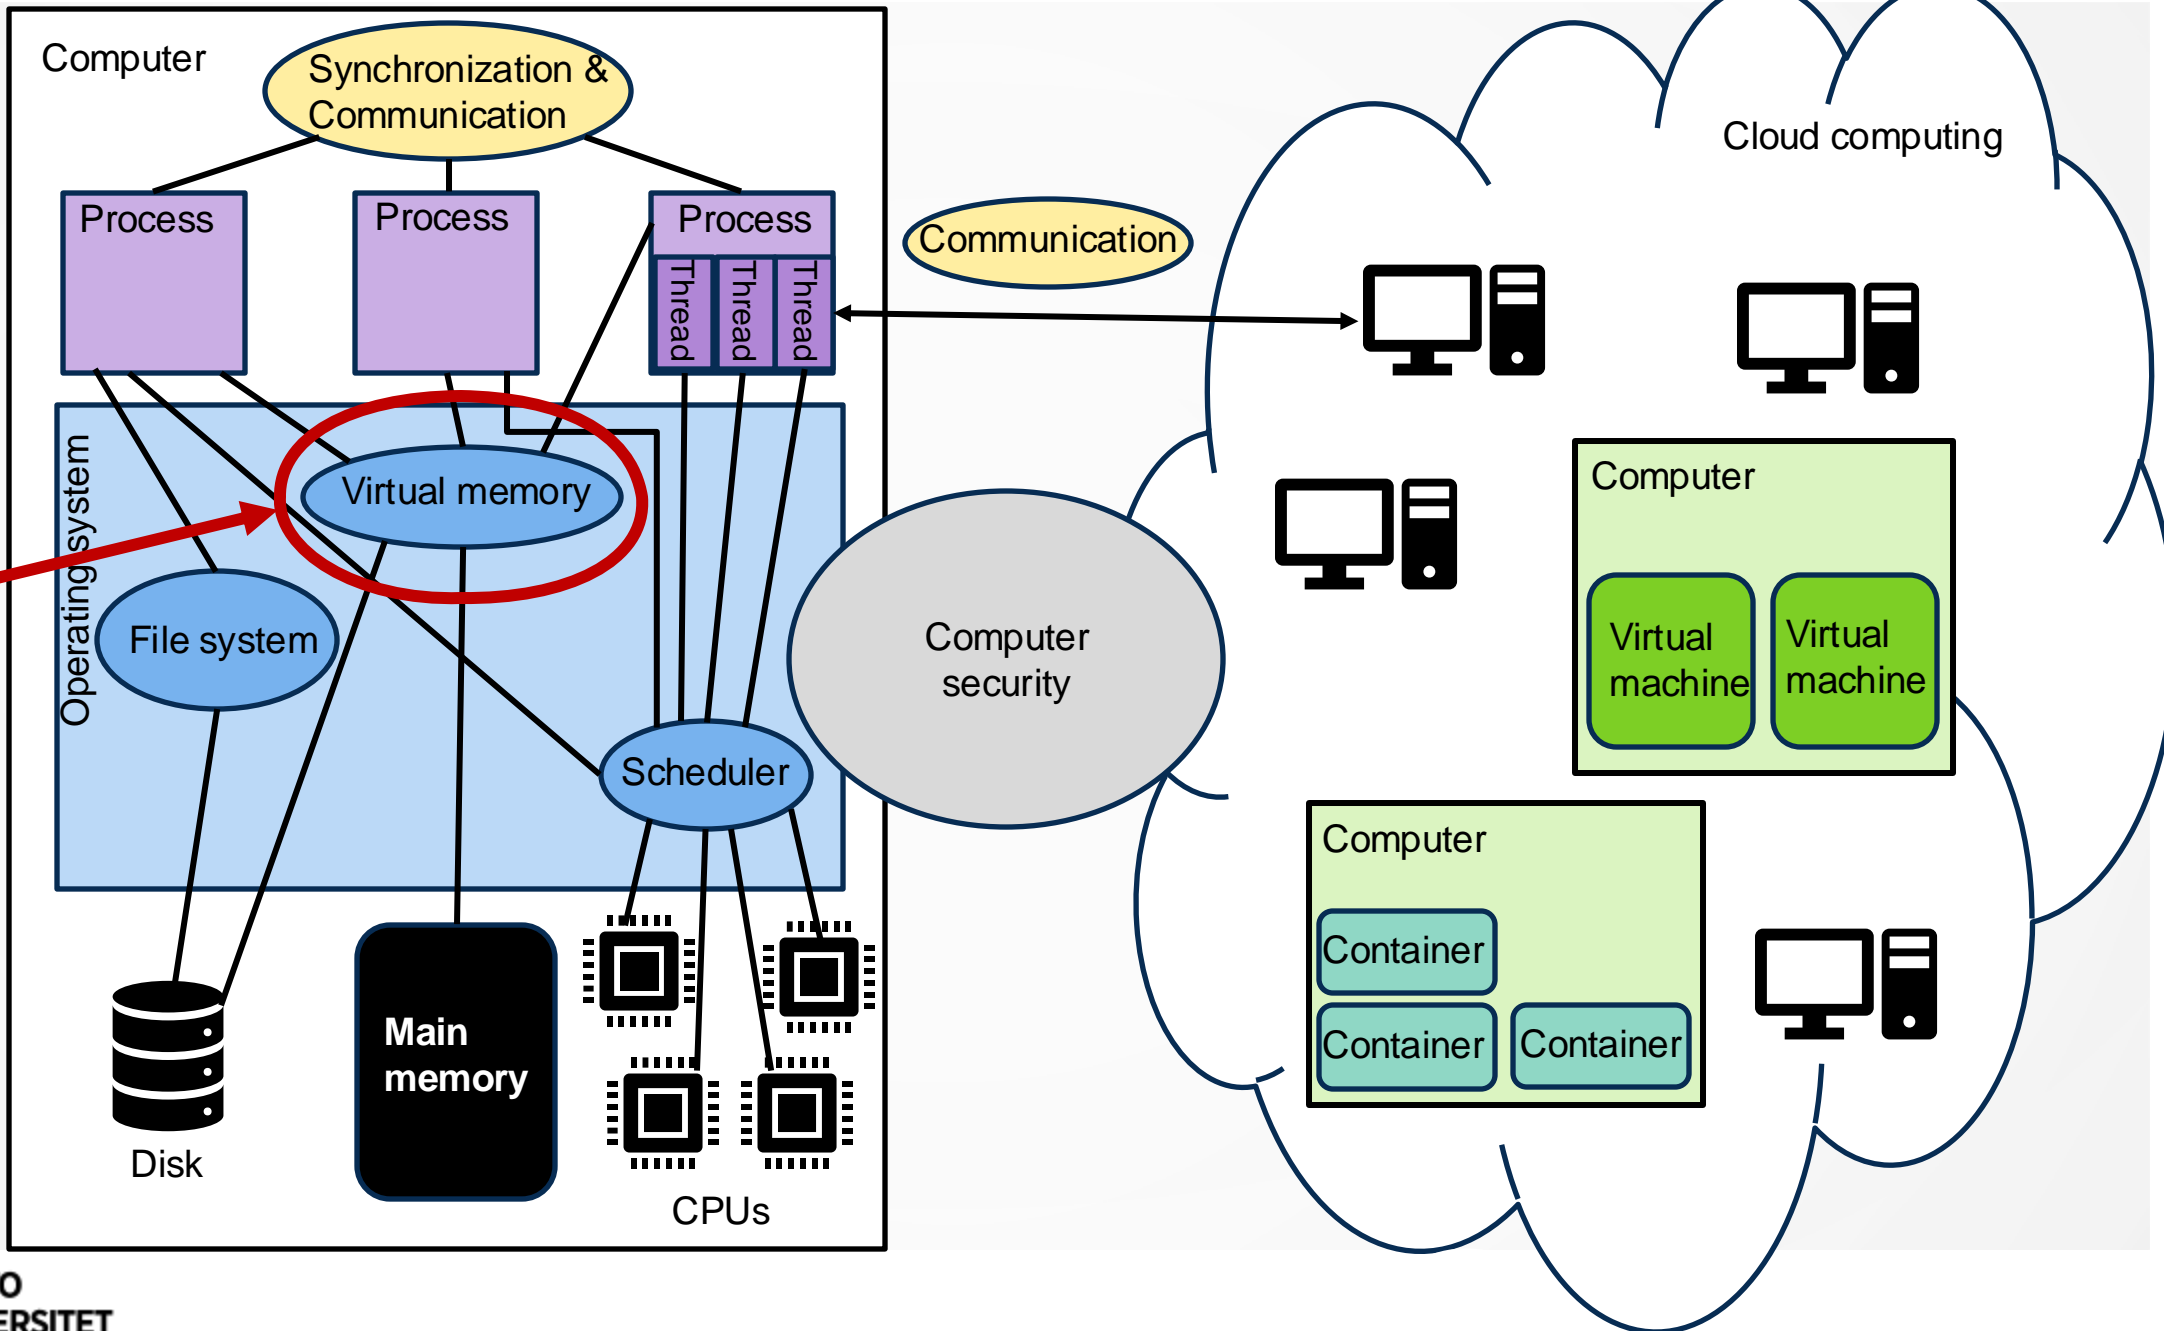


COMPUTING PLATFORMS

Virtual Memory:
Paging, segmentation, and address translation



Today's
topic





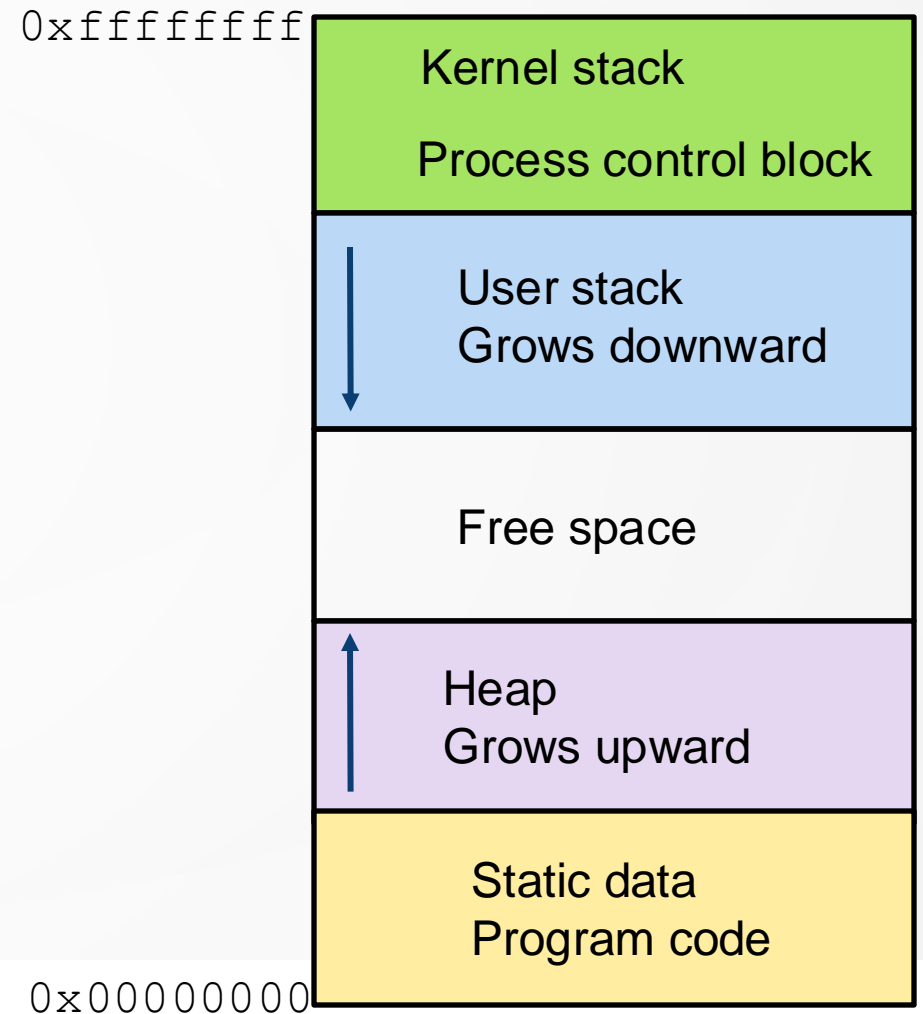
LEARNING OUTCOMES

- After today's lecture you
 - Know the requirements for memory management
 - Understand the concepts of segmentation and paging and can compare their advantages and disadvantages
 - Can describe hardware and OS structures that support segmentation and paging



REMINDER: HOW THE USER PROCESS SEES THE MEMORY

- Each process has its own linear address space and can reference addresses within its address space e.g. in range $0x00000000 - 0xffffffff$
- Most programs are organized into **modules**, some unmodifiable (read only, execute only) and some contain data that can be modified
 - Modules can be written and compiled independently, with references between modules resolved by the system at run time
 - Different degree of protection (read only, execute only) can be given to different modules
 - Some modules (library code) are shared between different processes





HOW THE OPERATING SYSTEM SEES MEMORY

- In a multiprogramming system, the main memory is divided to accommodate all processes
- The main memory may be smaller than the address space of a user process
- At least two levels of memory (in real systems more): **main memory** and **secondary memory**
 - Main memory provides fast access at relatively high cost
 - Main memory is volatile (not permanent storage)
 - Secondary memory is slower and cheaper but (usually) not volatile
 - Secondary memory of large capacity used for long-term storage of data and programs, while smaller main memory holds programs and data currently in use

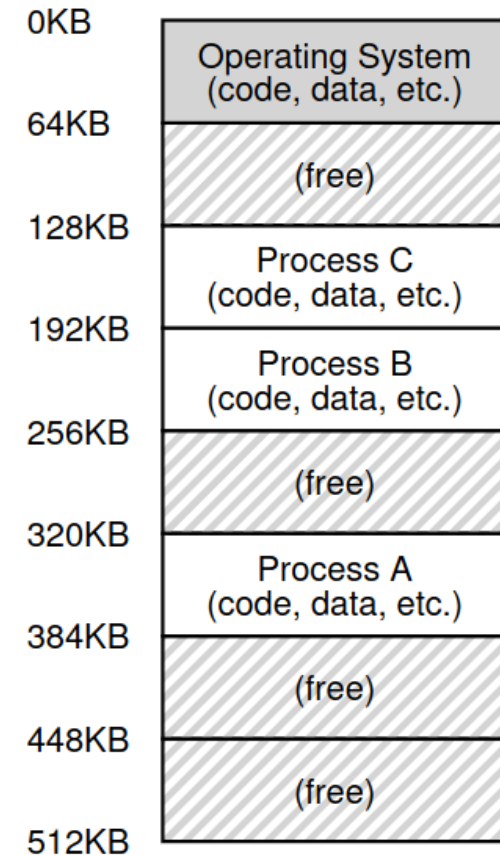


Figure 13.2: Three Processes: Sharing Memory



HOW TO ACCOMMODATE BOTH USER PROCESS'S AND OS'S POINT OF VIEW?

- User processes use **virtual addresses** to reference the process image
- **Address translation** is used to translate **virtual addresses** referencing the address space of the process to **physical addresses** referencing physical memory locations
- Both hardware and OS support needed to implement memory management!

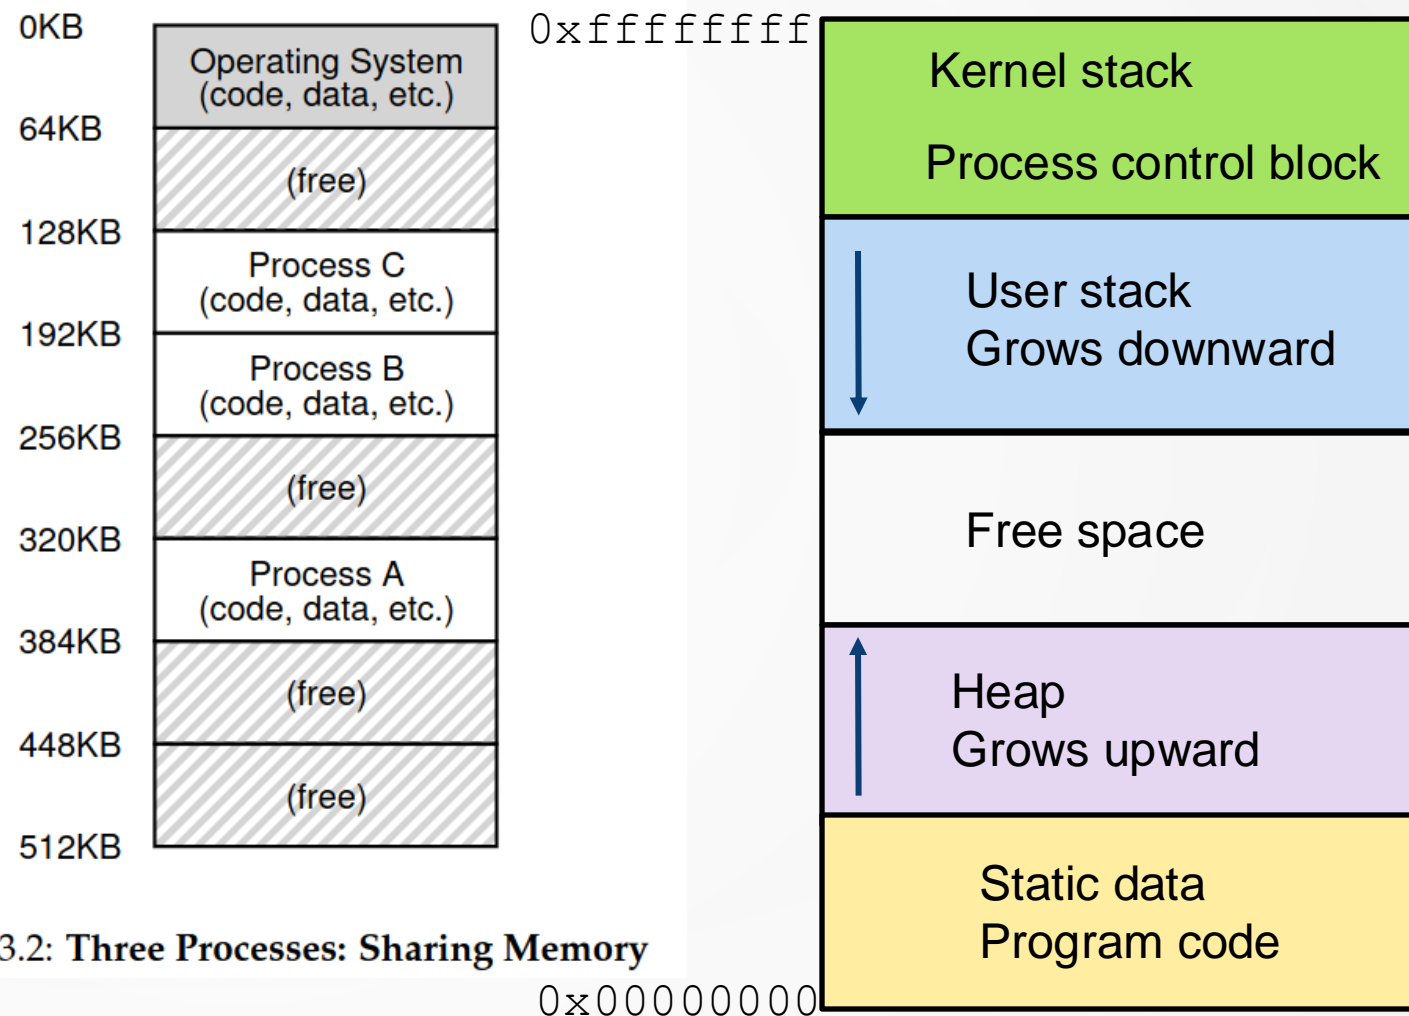


Figure 13.2: Three Processes: Sharing Memory



MEMORY MANAGEMENT REQUIREMENTS

- **Relocation:** the process image can be loaded to different locations of the physical memory
 - Invisible to the running program: Programmer should not need to know where in physical memory code, variables etc. will be placed
 - Enables swapping processes in and out of memory
- **Efficiency**
 - Programs should not run slower due to memory management
 - Structures supporting memory management should not take too much space
- **Protection & isolation**
 - Process should not be able to access or modify the address space of another process
 - Protect process from itself: some parts of address space should be execute-only (code), read-only, ...



RELOCATION AND ADDRESS TRANSLATION

- Simple example of address translation used in some early systems
- **Virtual address** (also **logical address**)
 - Reference to a memory location independent of the current assignment of data to memory
- **Relative address** (type of virtual address)
 - Address expressed as location relative to some known point (e.g. **base register**)
 - **Bounds register** gives limit of valid addresses
- **Physical address**
 - Actual location of data in main memory

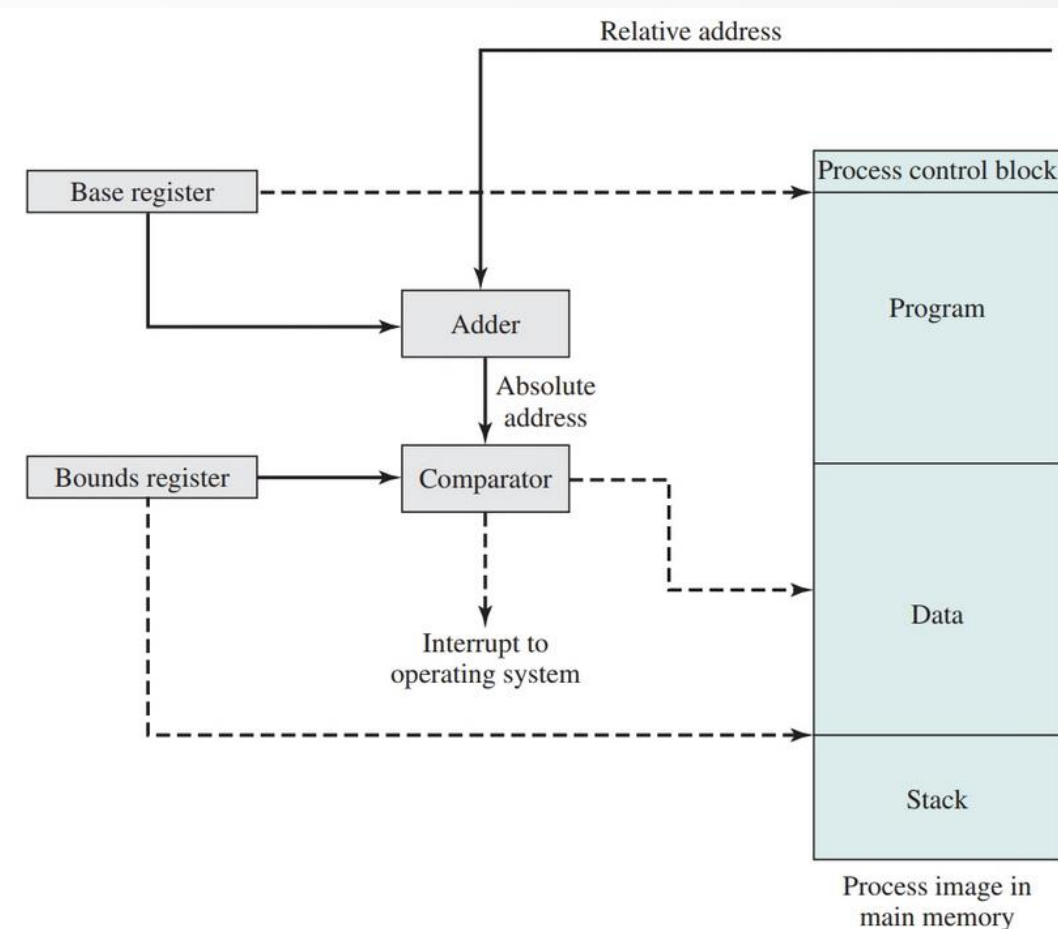


Figure 7.8 Hardware Support for Relocation



RELOCATION AND ADDRESS TRANSLATION: EXAMPLE

- Base register: 0001 0000 0000 0000 = 4096
- Bounds register: 0001 1000 0000 0000 = 6144
- Relative address: 0000 0011 0110 0111 = 439

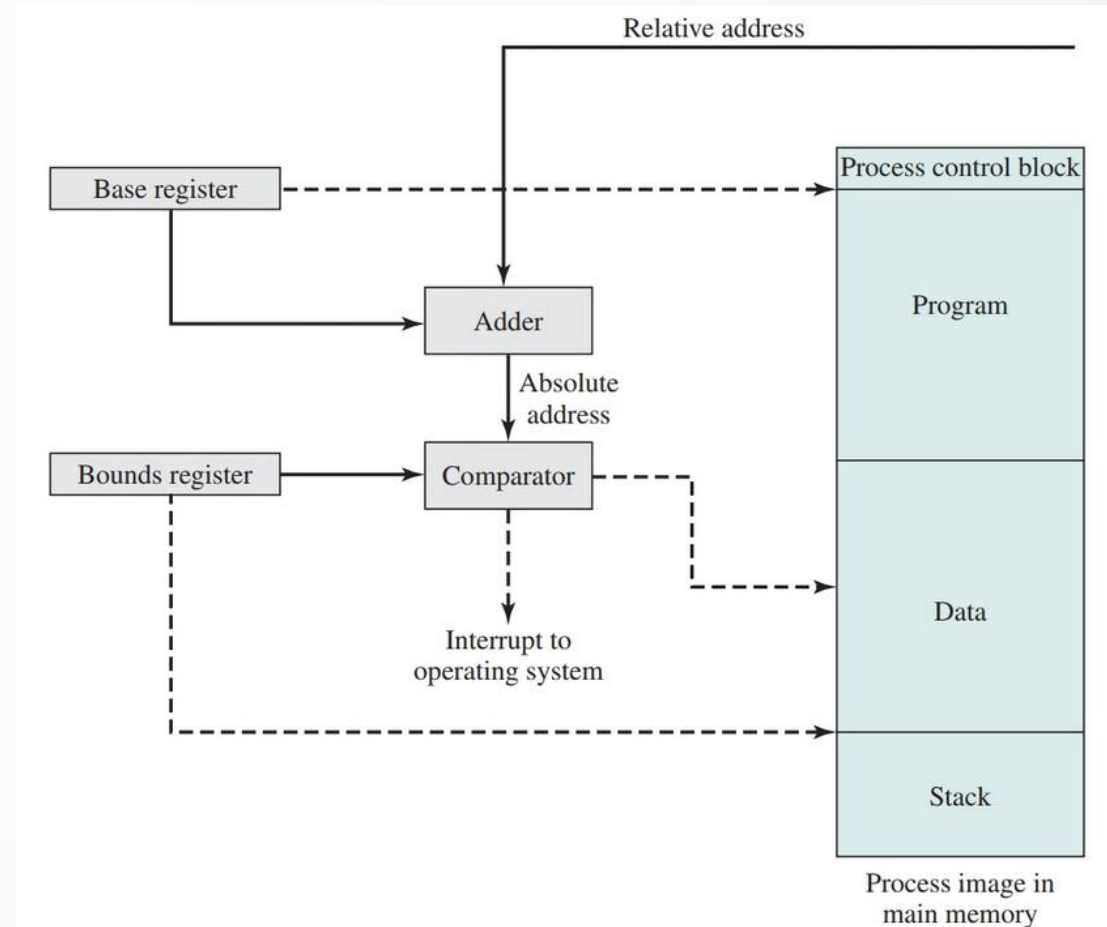


Figure 7.8 Hardware Support for Relocation



RELOCATION AND ADDRESS TRANSLATION: EXAMPLE

- Base register: 0001 0000 0000 0000 = 4096
- Bounds register: 0001 1000 0000 0000 = 6144
- Relative address: 0000 0011 0110 0111 = 439
- Physical address: Base + Relative address
= 0001 0000 0000 0000 + 0000 0011 0110 0111
= 0001 0011 0110 0111
= 4535
- Physical address is less than value in Bounds register -> Address is valid

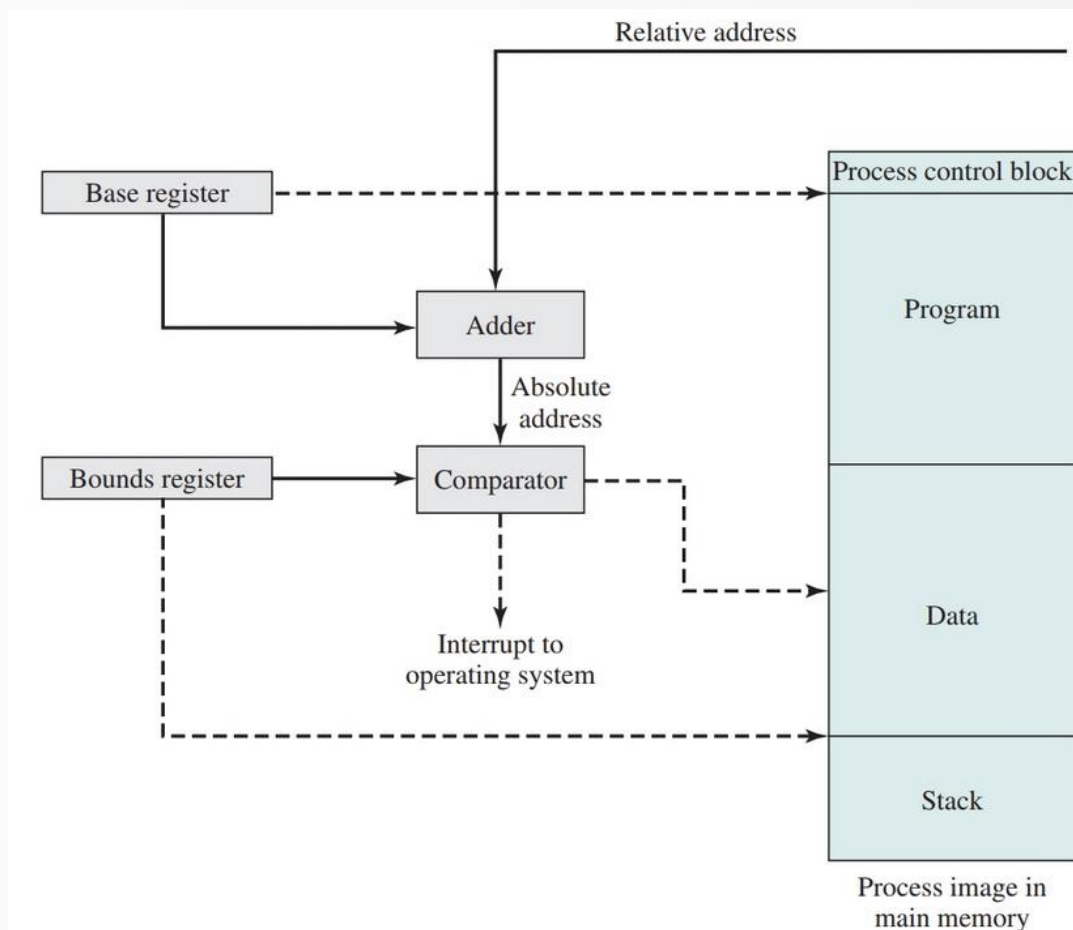


Figure 7.8 Hardware Support for Relocation



RELATIVE ADDRESSES

- Allows processes to be relocated (swapped in and out of memory)
- Process image is handled as one continuous block of memory
 - Need to find a continuous free block in main memory to accommodate the process
 - Cannot set different permissions (execute-only, read-only) for different parts of the process image
 - Does not provide a mechanism to share modules between processes
- Provides protection from other processes: Each process's addresses are contained within the values of Base and Bounds registers

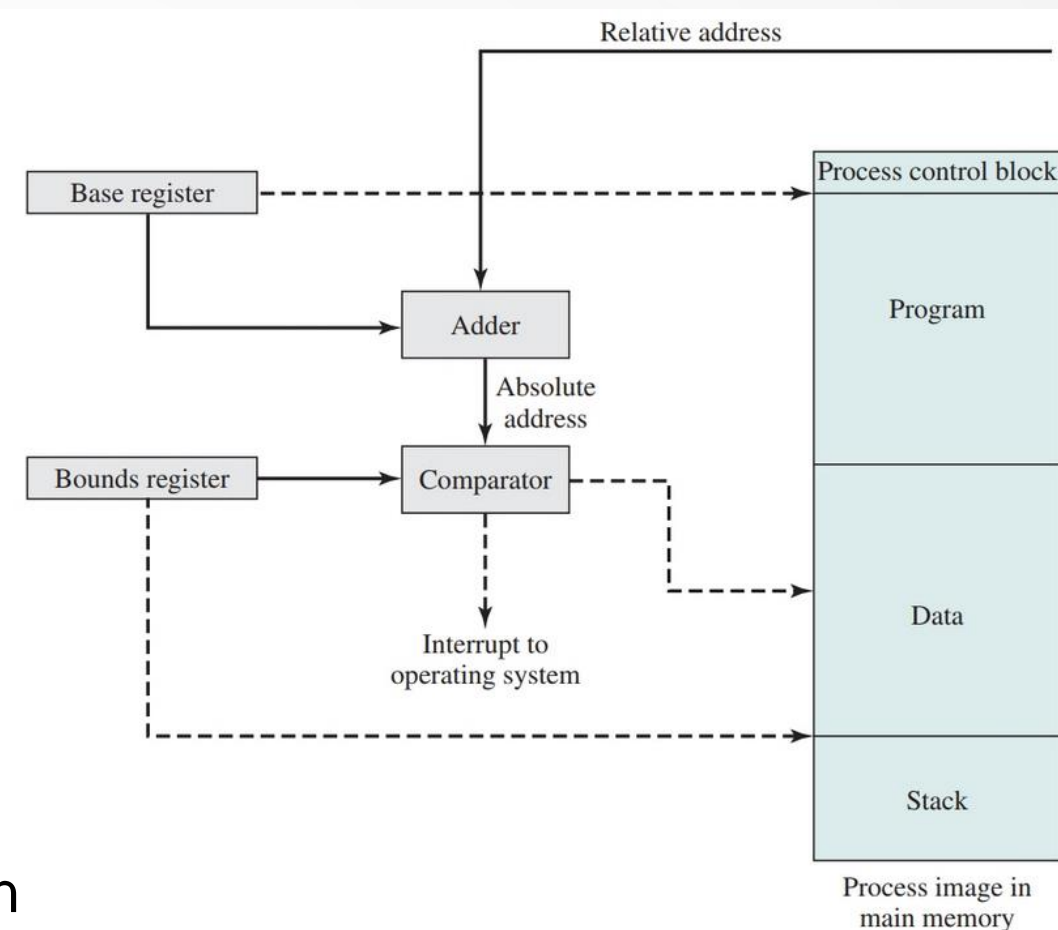


Figure 7.8 Hardware Support for Relocation

Figure from [Stallings, Operating systems: Internals and design principles, 9th ed]

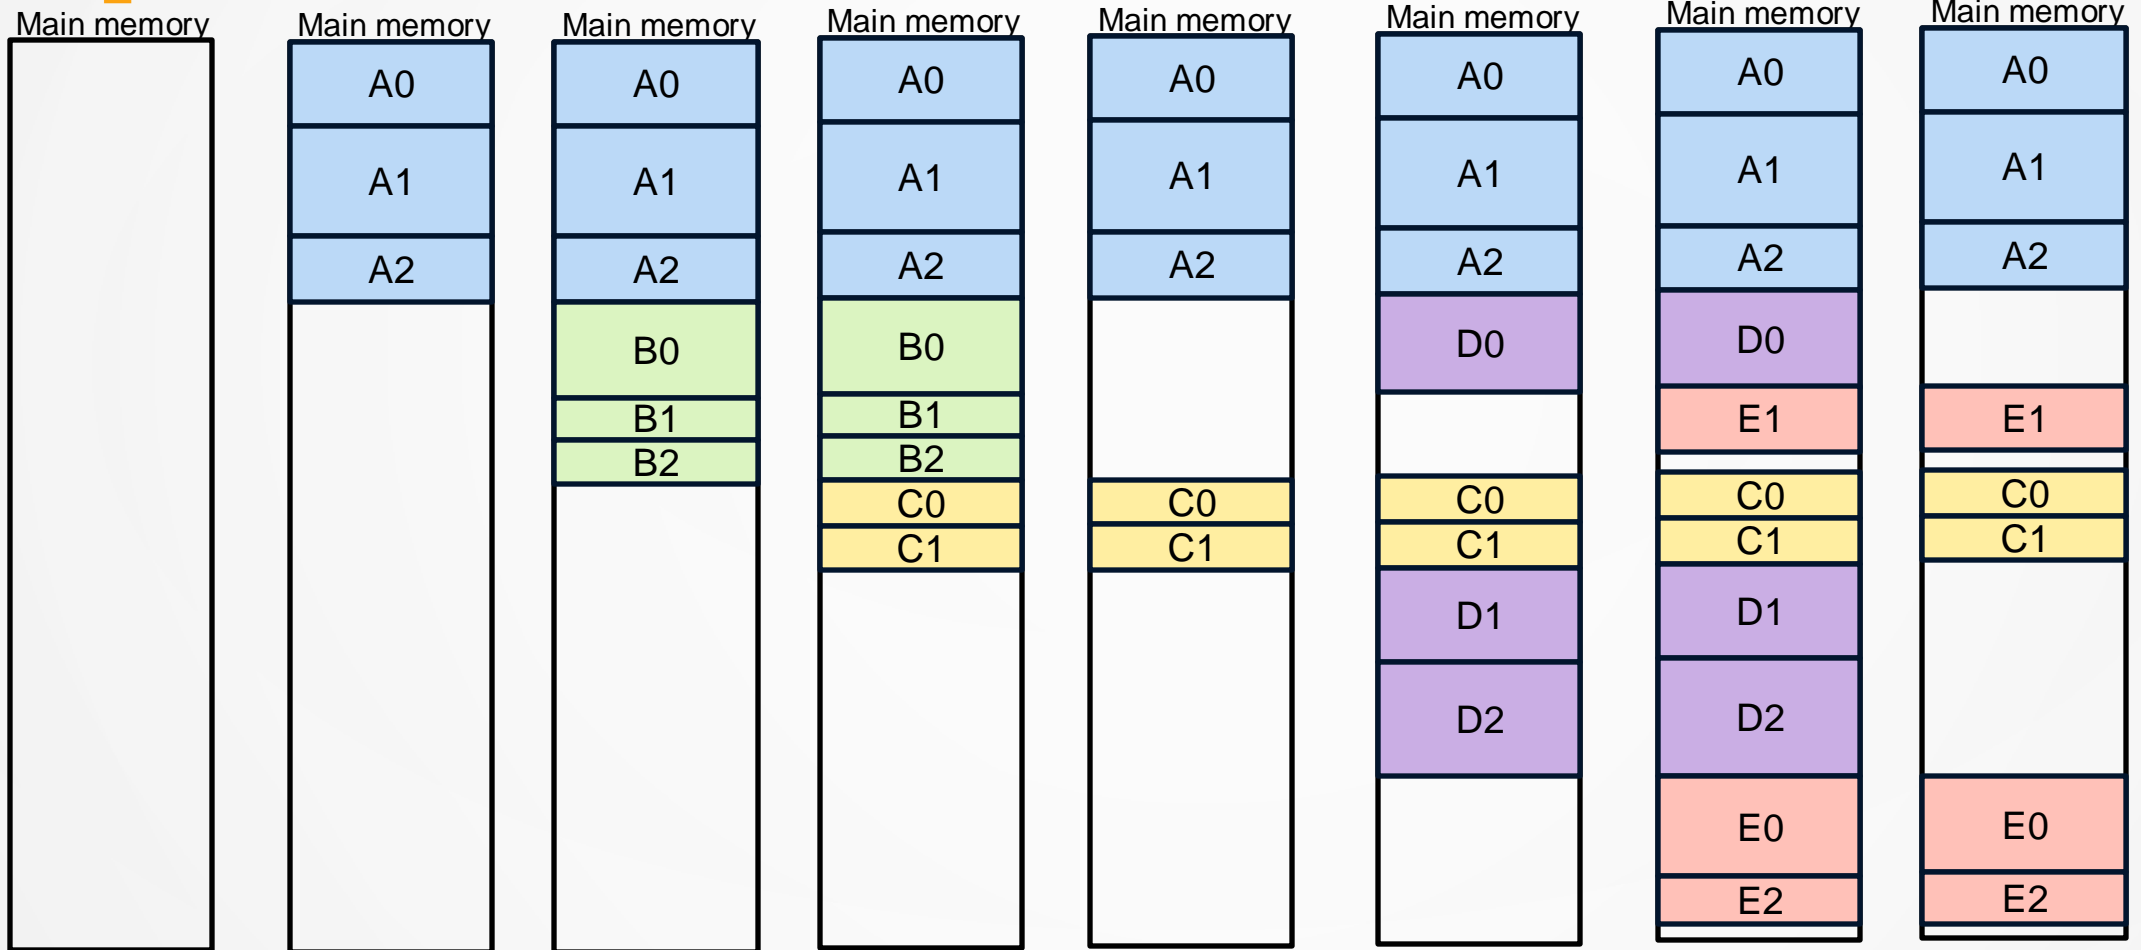


SEGMENTATION

- Program can be divided into **segments**
 - **Segment** is a variable length block of data (there is a maximum length)
- Virtual addresses consist of two parts
 - Segment number
 - Offset
- Visible to the programmer
- Typically code and data are assigned to different segments
 - Different permissions (execute-only, read-only) can be set for different segments
- Program or data can be broken down to multiple segments
 - Programmer must be aware of maximum segment size

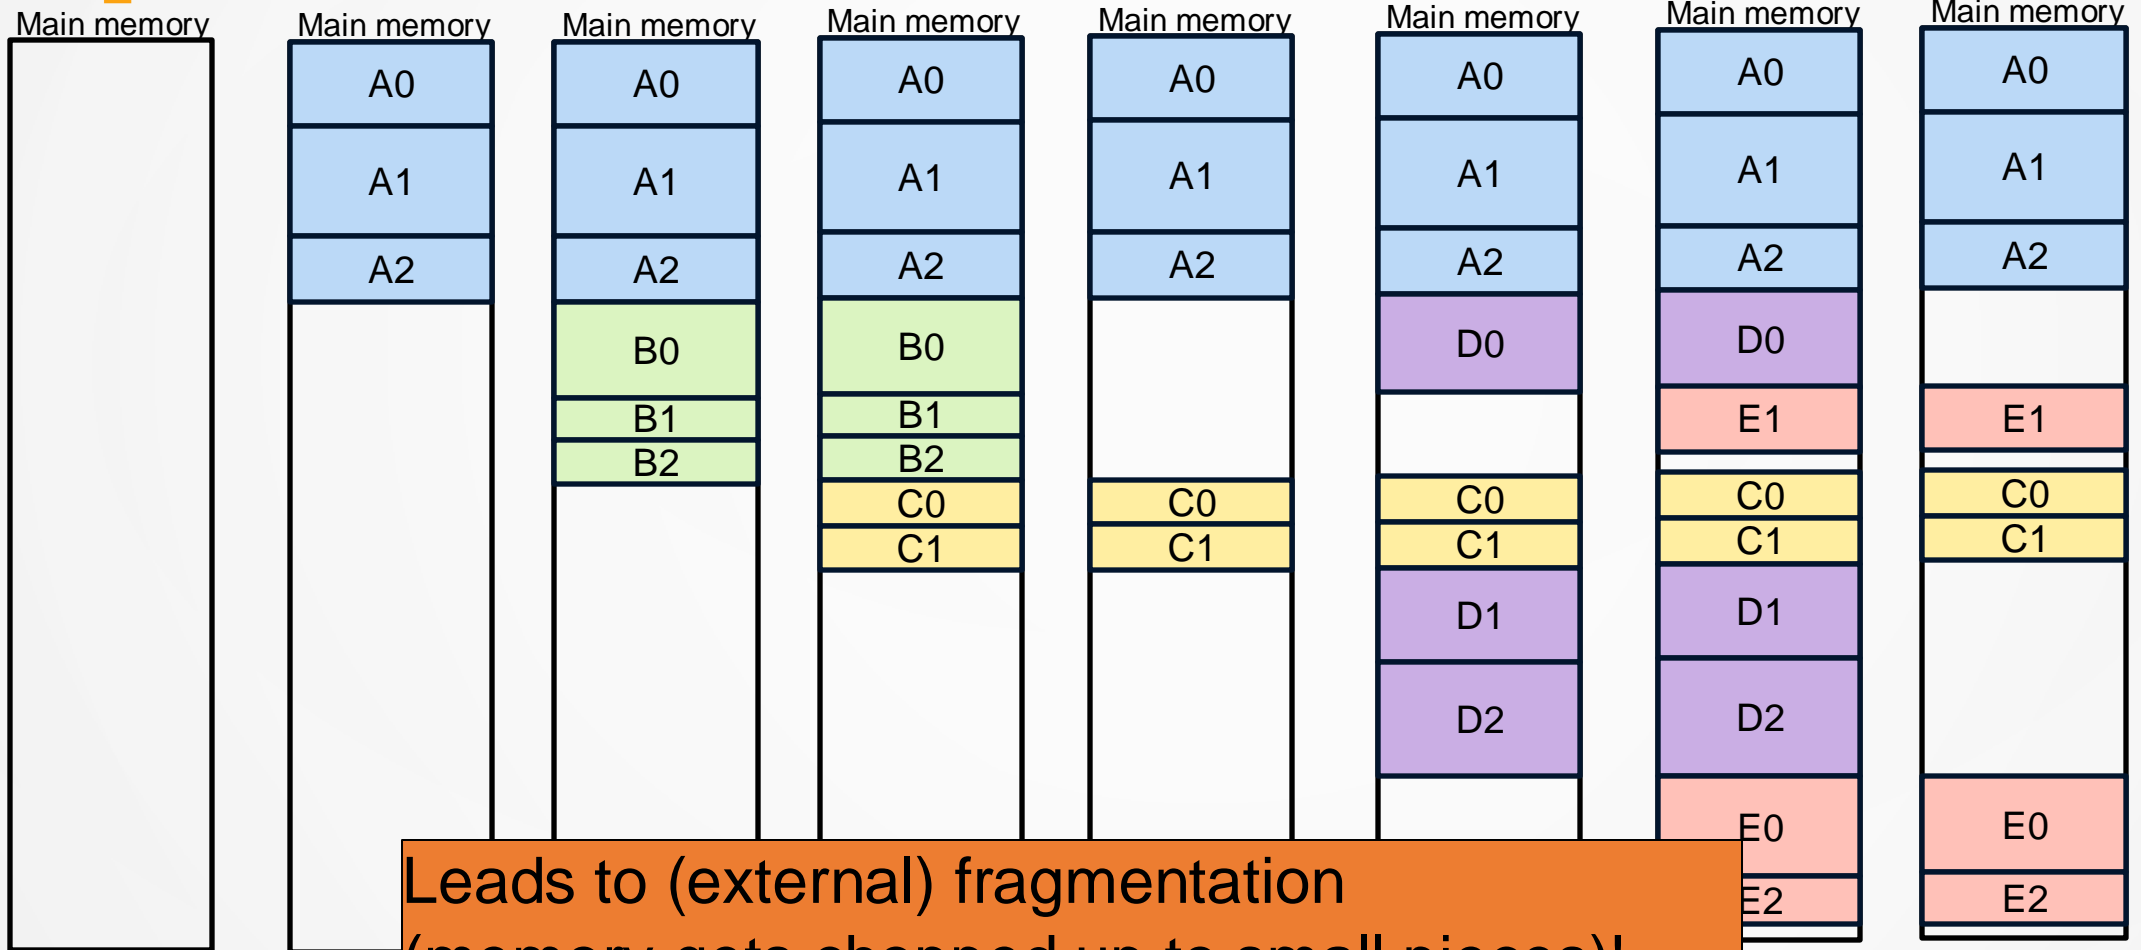


SEGMENTATION EXAMPLE





SEGMENTATION EXAMPLE



Leads to (external) fragmentation
(memory gets chopped up to small pieces)!



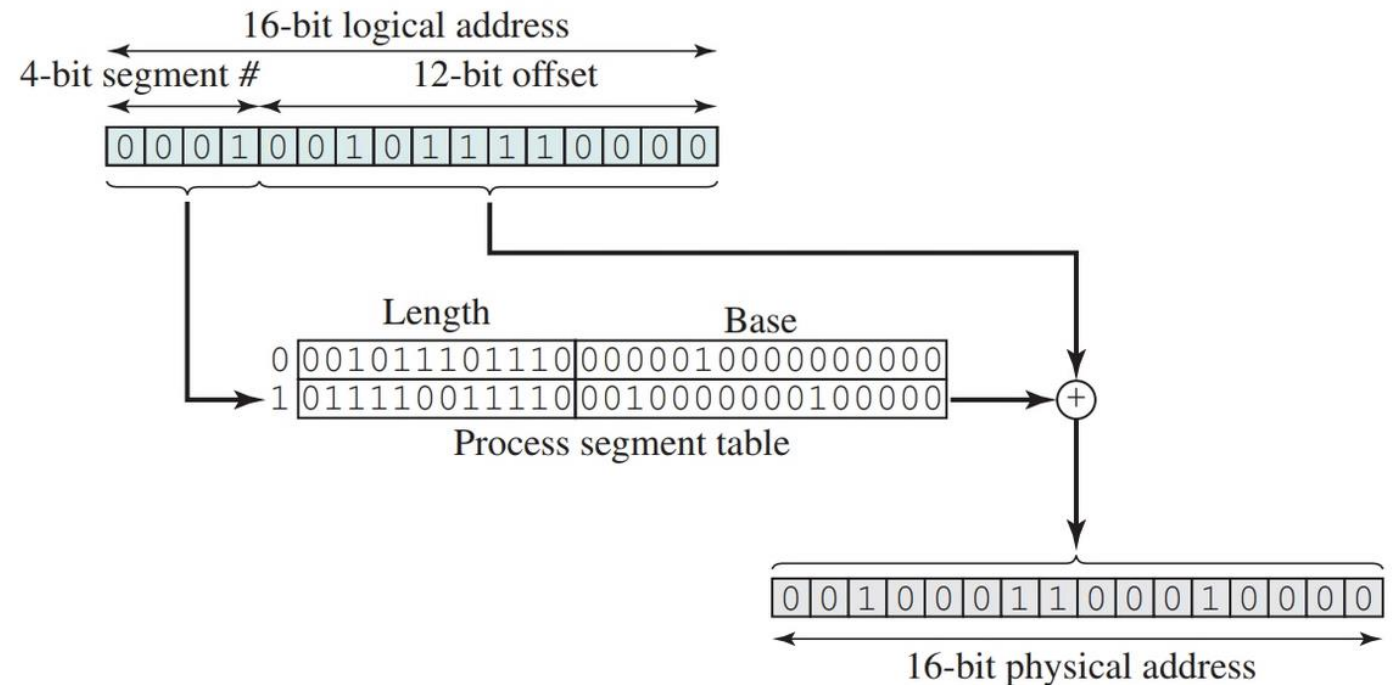
SEGMENTATION: ADDRESS TRANSLATION

- Due to unequal segment sizes, **there is no simple relationship between virtual and physical addresses**
- Consider a virtual address of $n+m$ bits, where the leftmost n bits give the segment number and the rightmost m bits give the offset.
- The steps of address translation
 1. **Extract the segment number** as the n leftmost bits of the virtual address
 2. Use the segment number as an index to the process segment table to find the starting physical address of the segment
 3. **Compare the offset** (rightmost m bits of the virtual address) to the length of the segment. If $\text{offset} \geq \text{segment length}$, **address is invalid**.
 4. Physical address is **sum of the starting physical address of the segment and the offset**



SEGMENTATION: ADDRESS TRANSLATION EXAMPLE

- Virtual address:
0001 0010 1111 0000



(b) Segmentation

Figure 7.12 Examples of Logical-to-Physical Address Translation



SEGMENTATION: ADDRESS TRANSLATION EXAMPLE

- Virtual address:
0001 0010 1111 0000
 - Segment 0001=1
 - Offset 0010 1111 0000 = 752
- According to process segment table, offset is smaller than segment length -> address is valid
- According to process segment table, this segment starts at physical address 0010 0000 0010 0000
- Physical address is then
$$\begin{array}{r} 0010\ 0000\ 0010\ 0000 \\ + \quad 0010\ 1111\ 0000 \\ \hline 0010\ 0011\ 0001\ 0000 \end{array}$$

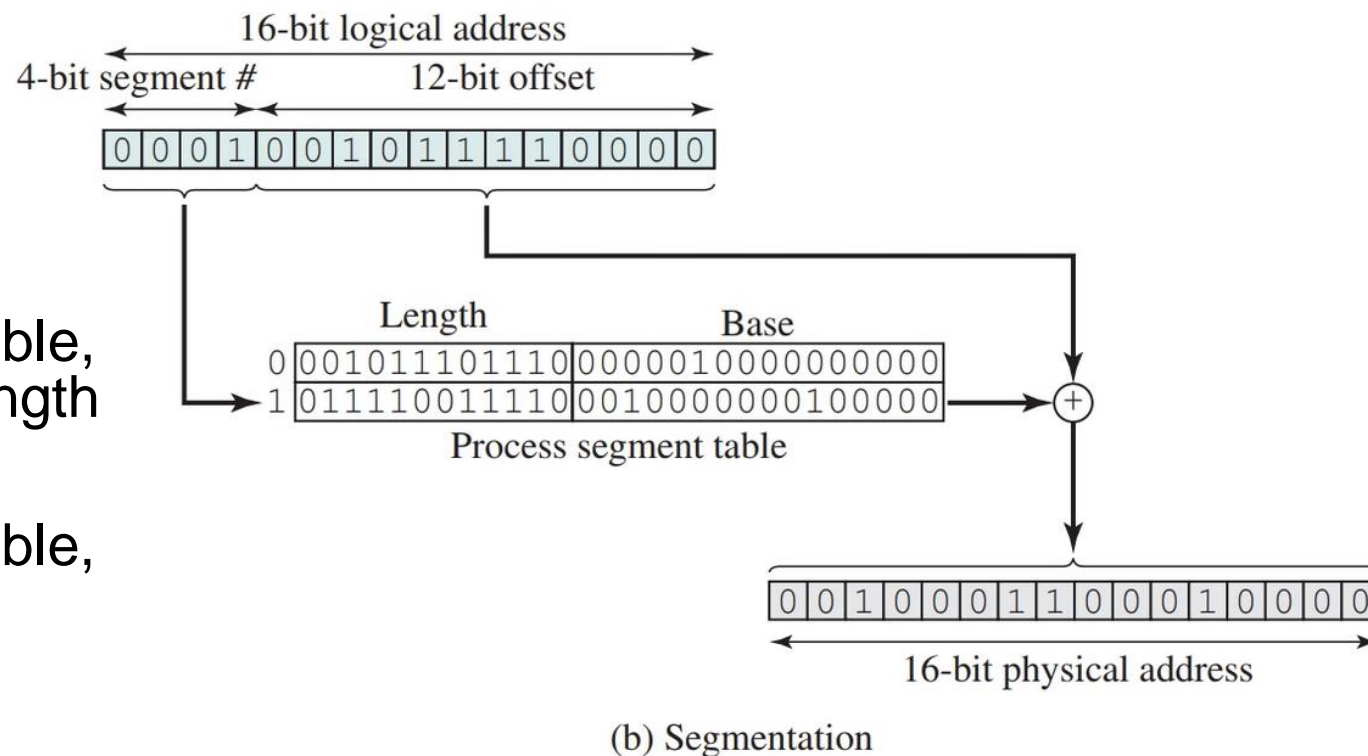


Figure 7.12 Examples of Logical-to-Physical Address Translation



SEGMENTATION: HARDWARE SUPPORT

- Memory management unit (MMU) has a **segment table pointer register**: points to the beginning of the current process's segment table
- Segment table has a structure dictated by MMU
- MMU does the address translation based on the segment table at runtime

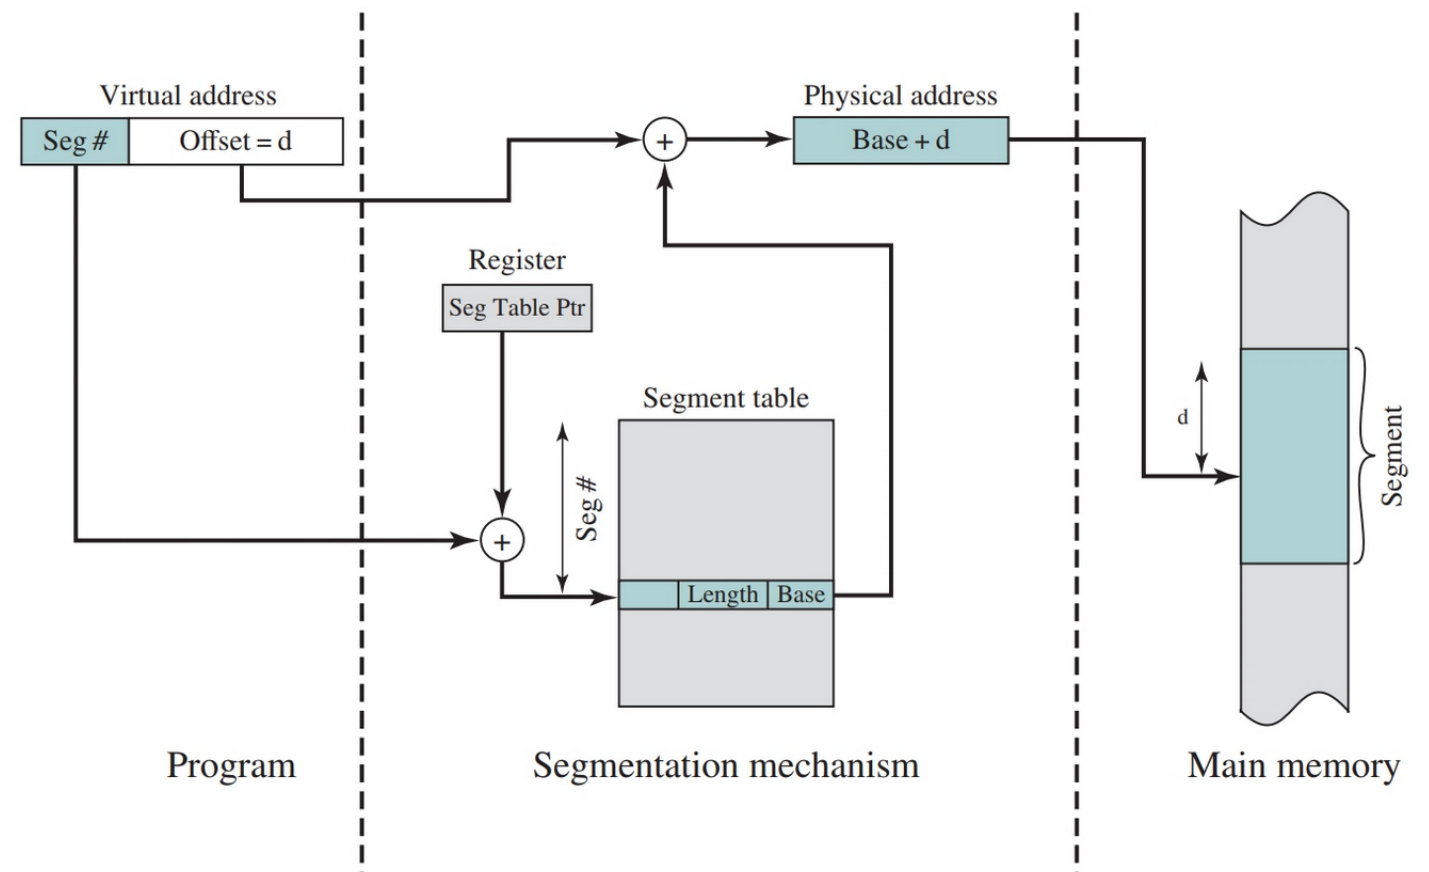


Figure 8.11 Address Translation in a Segmentation System



SEGMENTATION

- Allows processes to be relocated (swapped in and out of memory)
- Process image is divided into segments
 - Different permissions (execute-only, read-only) can be set for different segments
 - Allows sharing at segment level
 - Simplifies handling of growing data structures
 - Allows programs to be altered and recompiled independently
- Each segment is handled as one continuous block of memory
 - Need to find a continuous free block in main memory to accommodate the segment
 - Leads to **fragmentation**: free space gets chopped up to small pieces
- Provides protection from other processes: Each segment's addresses are contained within the values of base address and (base + length) as dictated in the segment table



PAGING

- Partition main memory into relatively small, **equal fixed-size** chunks
 - **Frames** = available chunks of memory
- Partition processes also into small, **fixed-size chunks** (same size!)
 - **Pages** = chunks of process (address space)
- **Page table**
 - OS maintains for each process
 - Contains frame location of each page in the process
 - Processor produces physical addresses (must know how to access the current process)
- Invisible to the programmer



PAGING: EXAMPLE

Frame number	Main memory	Frame number	Main memory	Frame number	Main memory	Frame number	Main memory	Frame number	Main memory	Frame number	Main memory
0		0	A0	0	A0	0	A0	0	A0	0	A0
1		1	A1	1	A1	1	A1	1	A1	1	A1
2		2	A2	2	A2	2	A2	2	A2	2	A2
3		3	A3	3	A3	3	A3	3	A3	3	A3
4		4		4	B0	4	B0	4		4	D0
5		5		5	B1	5	B1	5		5	D1
6		6		6	B2	6	B2	6		6	D2
7		7		7		7	C0	7	C0	7	C0
8		8		8		8	C1	8	C1	8	C1
9		9		9		9	C2	9	C2	9	C2
10		10		10		10	C3	10	C3	10	C3
11		11		11		11		11		11	D3
12		12		12		12		12		12	
13		13		13		13		13		13	



PAGING: EXAMPLE

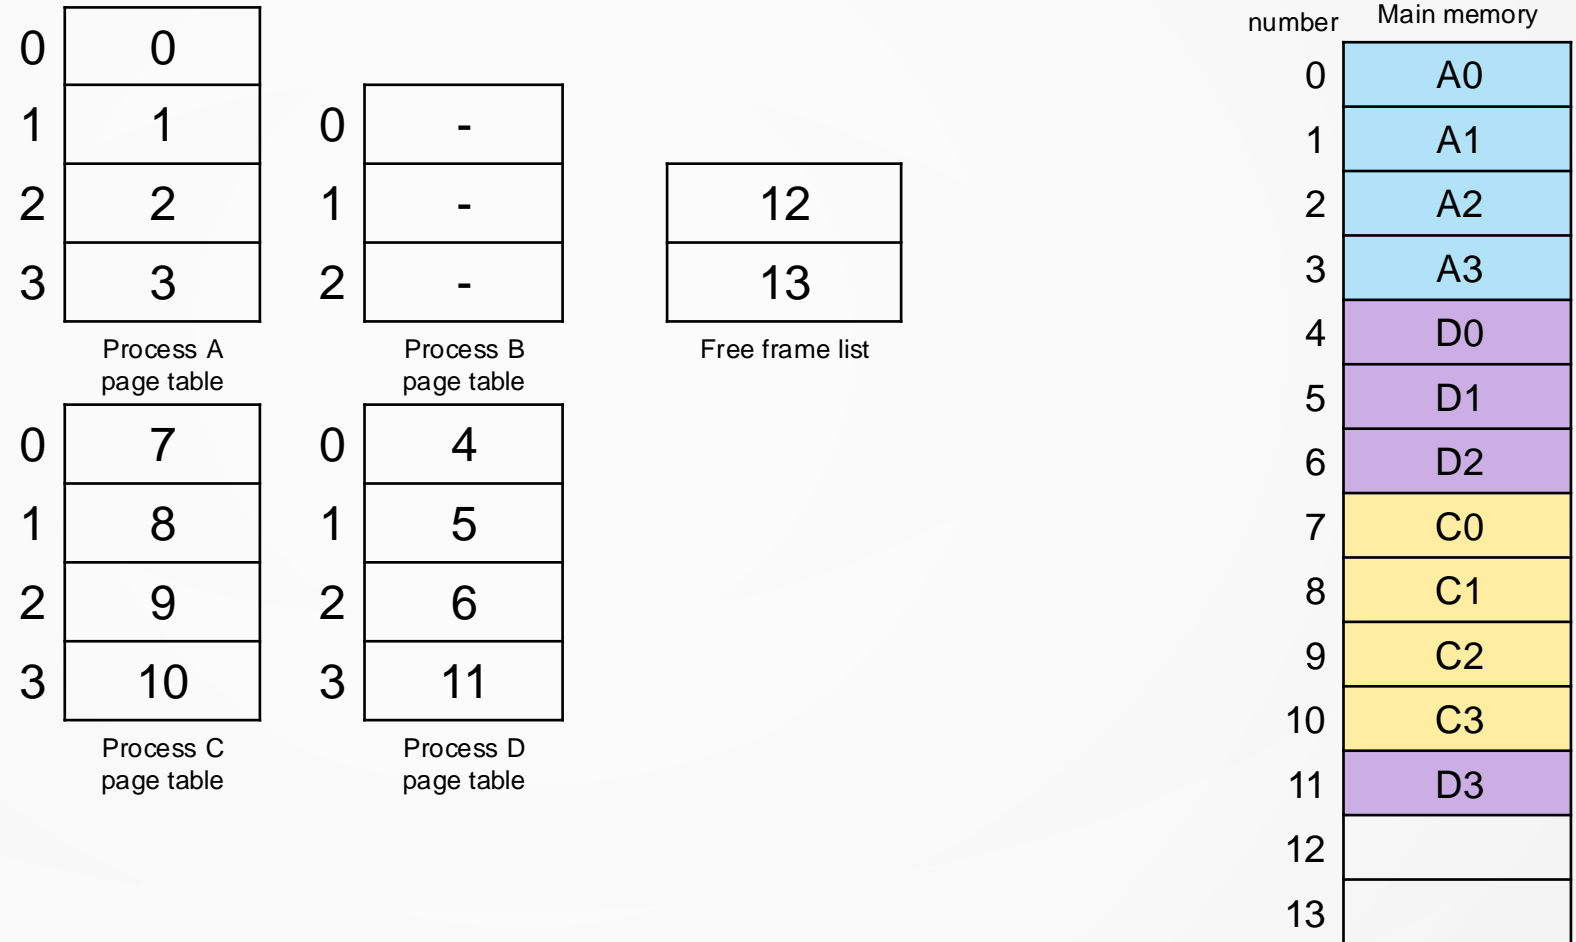
Frame number	Main memory	Frame number	Main memory	Frame number	Main memory	Frame number	Main memory	Frame number	Main memory
0		0	A0	0	A0	0	A0	0	A0
1		1	A1	1	A1	1	A1	1	A1
2		2	A2	2	A2	2	A2	2	A2
3		3	A3	3	A3	3	A3	3	A3
4		4		4	B0	4	B0	4	D0
5		5		5	B1	5	B1	5	D1
6		6		6	B2	6	B2	6	D2
7		7		7		7	C0	7	C0
8		8		8		8	C1	8	C1
9		9		9		9	C2	9	C2
10		10		10		10	C3	10	C3
11		11		11		11		11	D3
12		12		12		12		12	
13		13		13		13		13	

No external fragmentation!



PAGING: EXAMPLE – PAGE TABLES

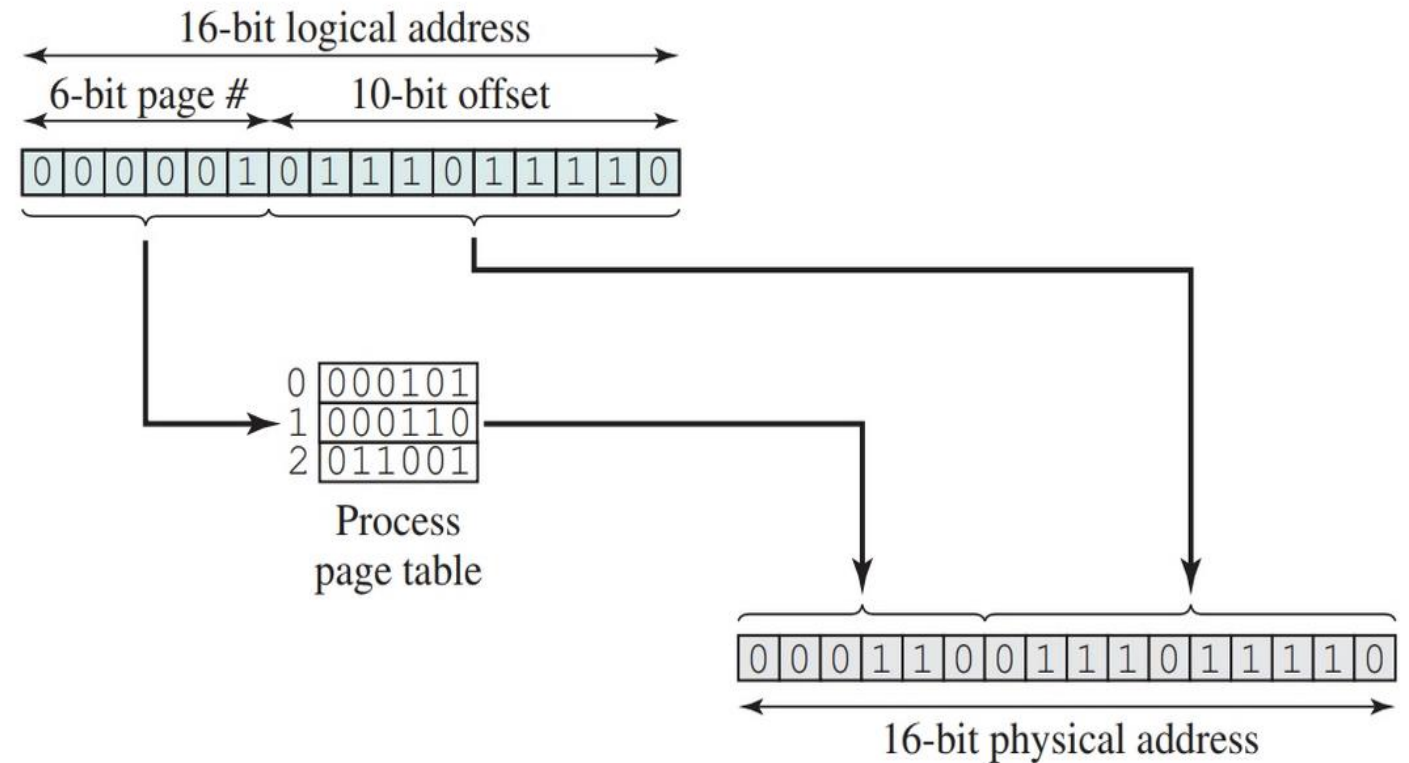
- Process page tables and the free space list for the example (last step)





PAGING: ADDRESS TRANSLATION EXAMPLE

- Virtual address
0000 0101 1101 1110



(a) Paging

Figure 7.12 Examples of Logical-to-Physical Address Translation



PAGING: ADDRESS TRANSLATION EXAMPLE

- Virtual address
0000 0101 1101 1110
 - Page number: 0000 01 = 1
 - Offset: 01 1101 1110 = 478
- According to the process page table page 1 resides in frame
0001 10=6
- The physical address is then the concatenation of frame number and offset:
0001 1001 1101 1110

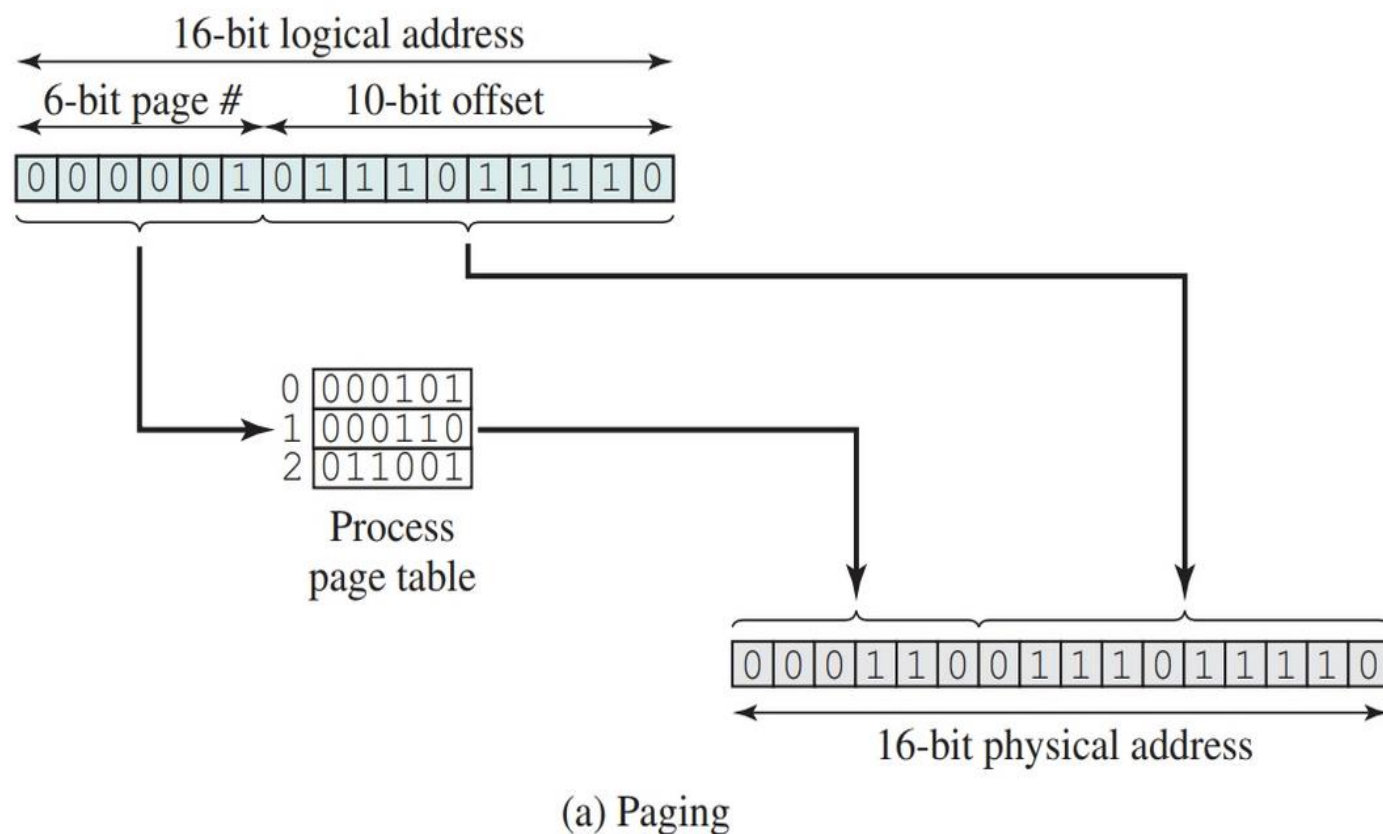


Figure 7.12 Examples of Logical-to-Physical Address Translation



PAGE SIZE: WHY A POWER OF 2?

- Relatively easy to **implement a function in hardware** to perform dynamic **address translation at run time**
- Consider address of $n+m$ bits, where the leftmost n bits are the page number and the rightmost m bits are the offset. Then:
 1. **Extract the page number** (leftmost n bits of the address)
 2. **Use page number as index** into the process page table to find frame number k
 3. **Physical address is $k*2^m + \text{offset}$ (this needs not be calculated; the bits of the frame number and the bits of the offset are just concatenated)**



PAGING: HARDWARE SUPPORT

- MMU has a **page table pointer register**: points to the beginning of the page table of the current process
- Page table structure is dictated by the MMU
- MMU does address translation at run time

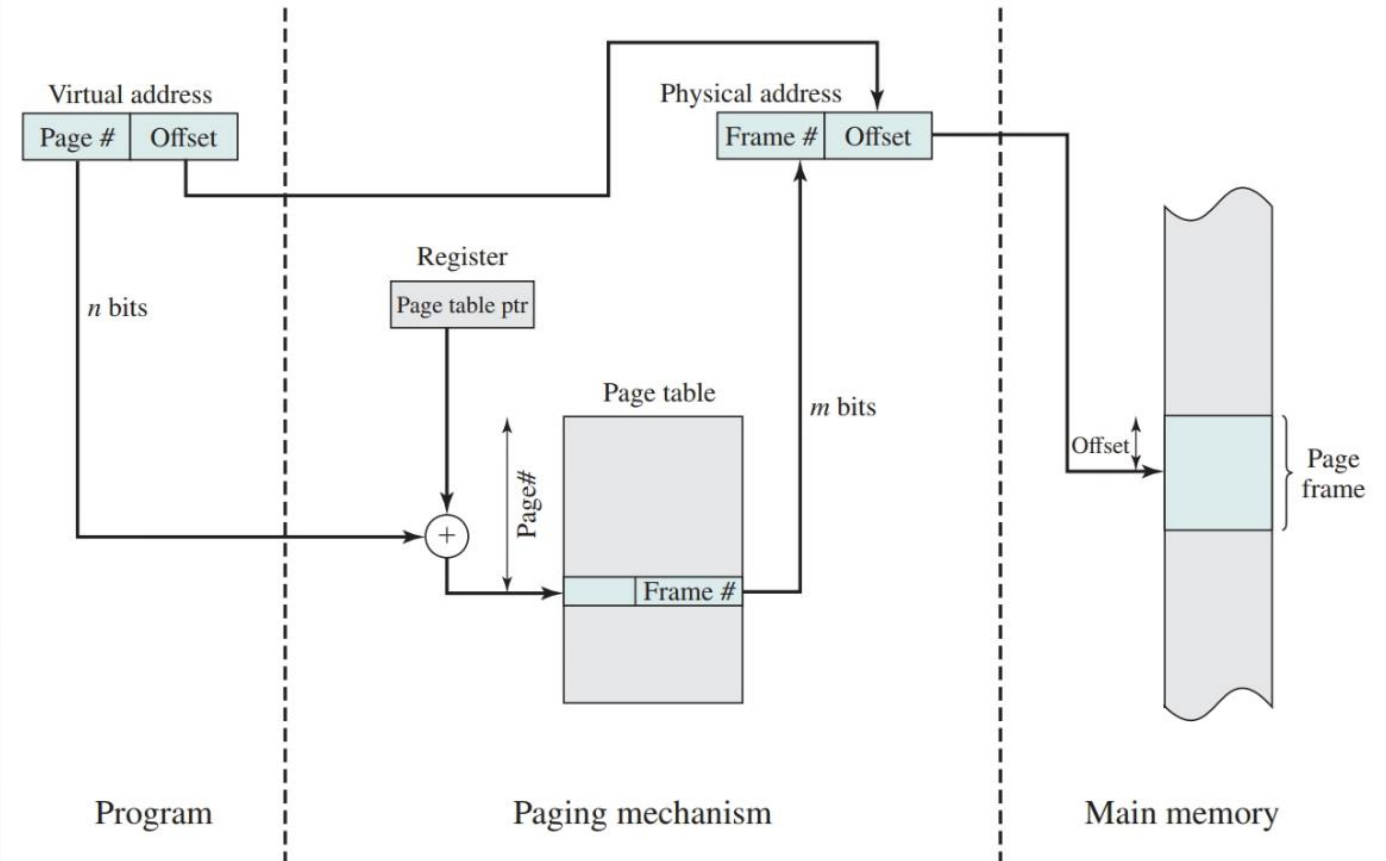


Figure 8.2 Address Translation in a Paging System



PAGING

- Allows processes to be relocated (swapped in and out of memory)
- Process image is divided into pages
 - Different permissions (execute-only, read-only) can be set for different pages
 - Allows sharing at page level
 - However, pages are not visible to the programmer!
- Pages and frames have a fixed size
 - Easy to manage free space
 - No **external fragmentation**
 - **Internal fragmentation:** wasted space due to a page not being fully used
- Provides protection from other processes: Each page is only visible to the process to which it belongs to



PAGE SIZE

- The smaller the page size, the less internal fragmentation
- However...
 - ... more pages required per process
 - ... more pages per process means larger page tables
 - ... physical characteristics of most secondary-memory devices favor large page size for more **efficient block transfer of data**
- Typically, page size is 0.5 KB to 8 KB



HOW BIG IS A PAGE TABLE?

- Consider 32-bit virtual addresses with 4 KB pages
 - 20-bit virtual page numbers and 12-bit offsets
- OS needs a page table entry for each possible virtual page
 - Assuming 4 bytes per entry, the page table would take $2^{20} * 4 \text{ B} = 4 \text{ MB}$
 - Assuming system has 100 processes, this means 400 MB for page tables only!
- What about 64-bit virtual addresses with 4 KB pages?
 - 52-bit virtual page numbers and 12-bit offsets
 - Assuming 8 bytes per entry, the page table would take $2^{52} * 8 \text{ B} = 32768 \text{ TB}$



TWO-LEVEL HIERARCHICAL PAGE TABLE

- In practice, even 5-level page tables exist in today's systems




Figure 8.3 A Two-Level Hierarchical Page Table



TWO-LEVEL HIERARCHICAL PAGE TABLE: ADDRESS TRANSLATION

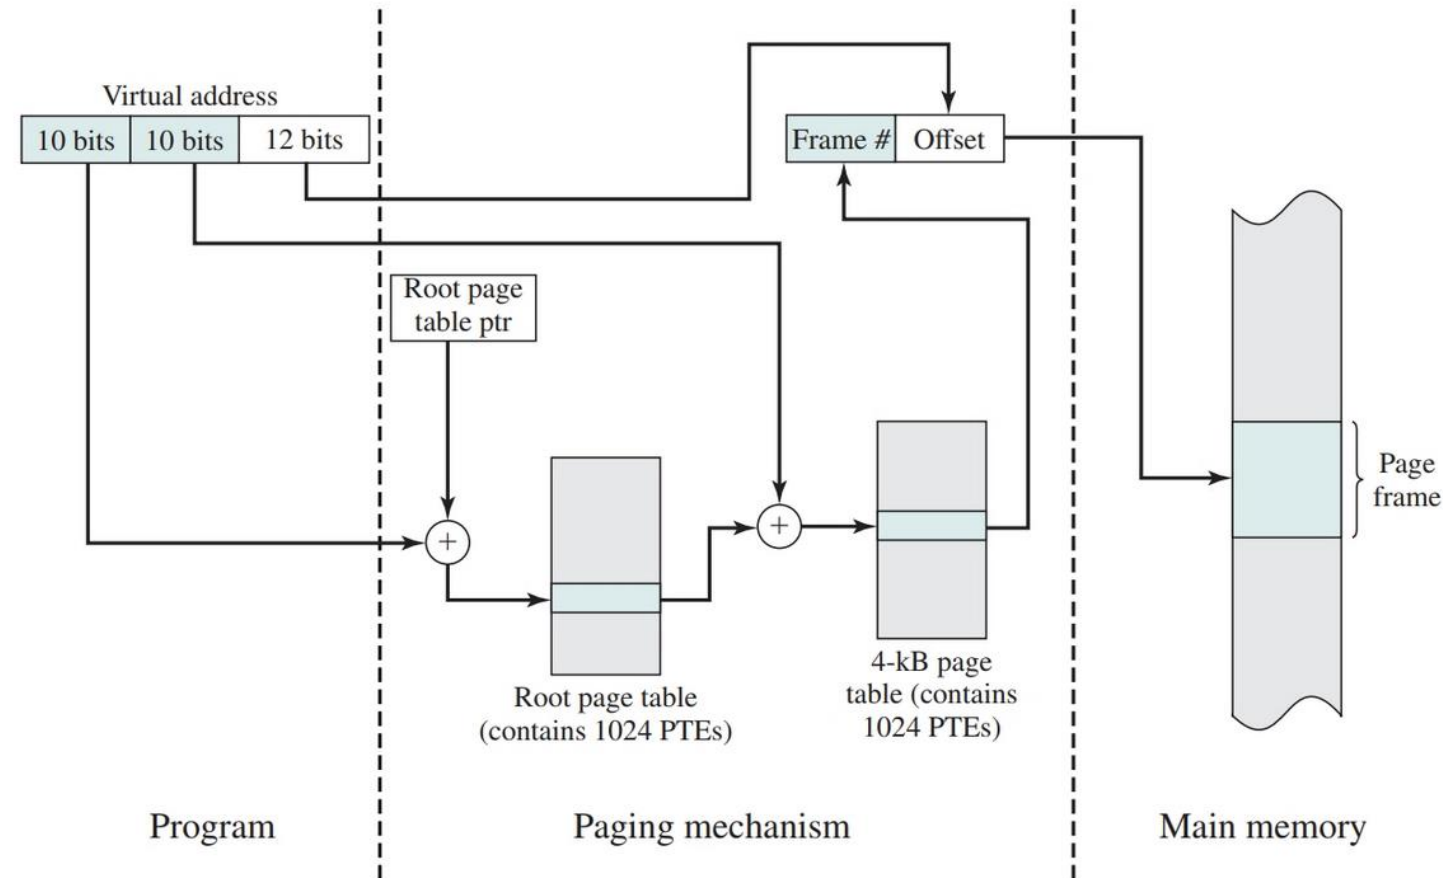


Figure 8.4 Address Translation in a Two-Level Paging System



TWO-LEVEL HIERARCHICAL PAGE TABLE: ADDRESS TRANSLATION EXAMPLE

- Virtual address:
0000 0001 1000 0000 1000 0101 0111 0010

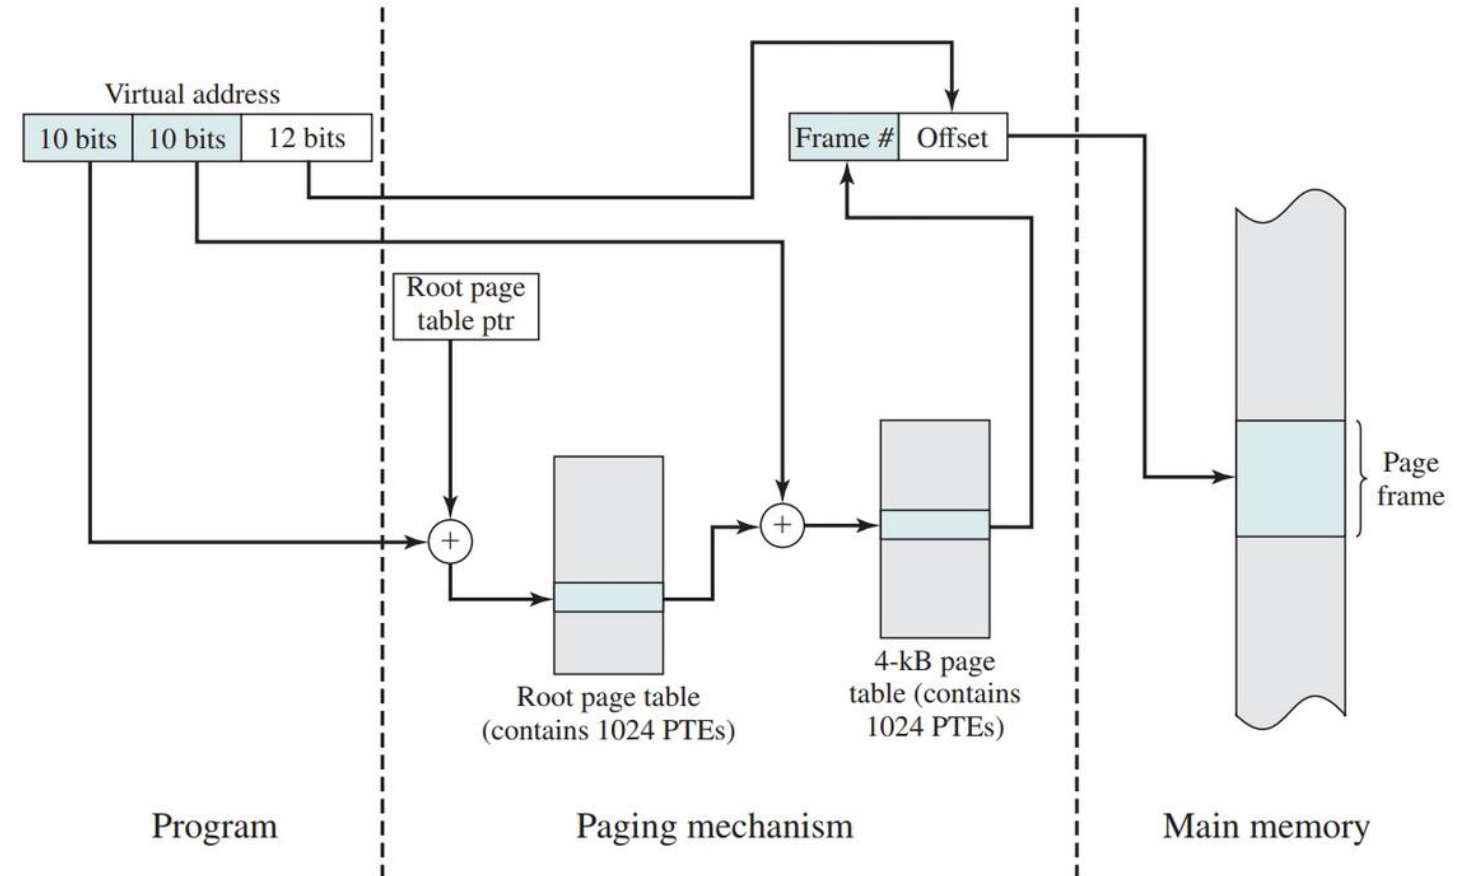


Figure 8.4 Address Translation in a Two-Level Paging System



TWO-LEVEL HIERARCHICAL PAGE TABLE: ADDRESS TRANSLATION EXAMPLE

- Virtual address:
0000 0001 1000 0000 1000 0101 0111 0010
 - Index to root page table: 00 0000 0110 = 6
 - Index to 4-kB page table: 00 0000 1000 = 8
 - Offset: 0101 0111 0010
- 6th entry in root page table gives the address of the 4-kB page table
- 8th entry in 4-kB page table gives the frame # (assume it is 3 = 0x0000 0000 0000 0011)
- Physical address:
0x0000 0000 0000 0011 0101 0111 0010

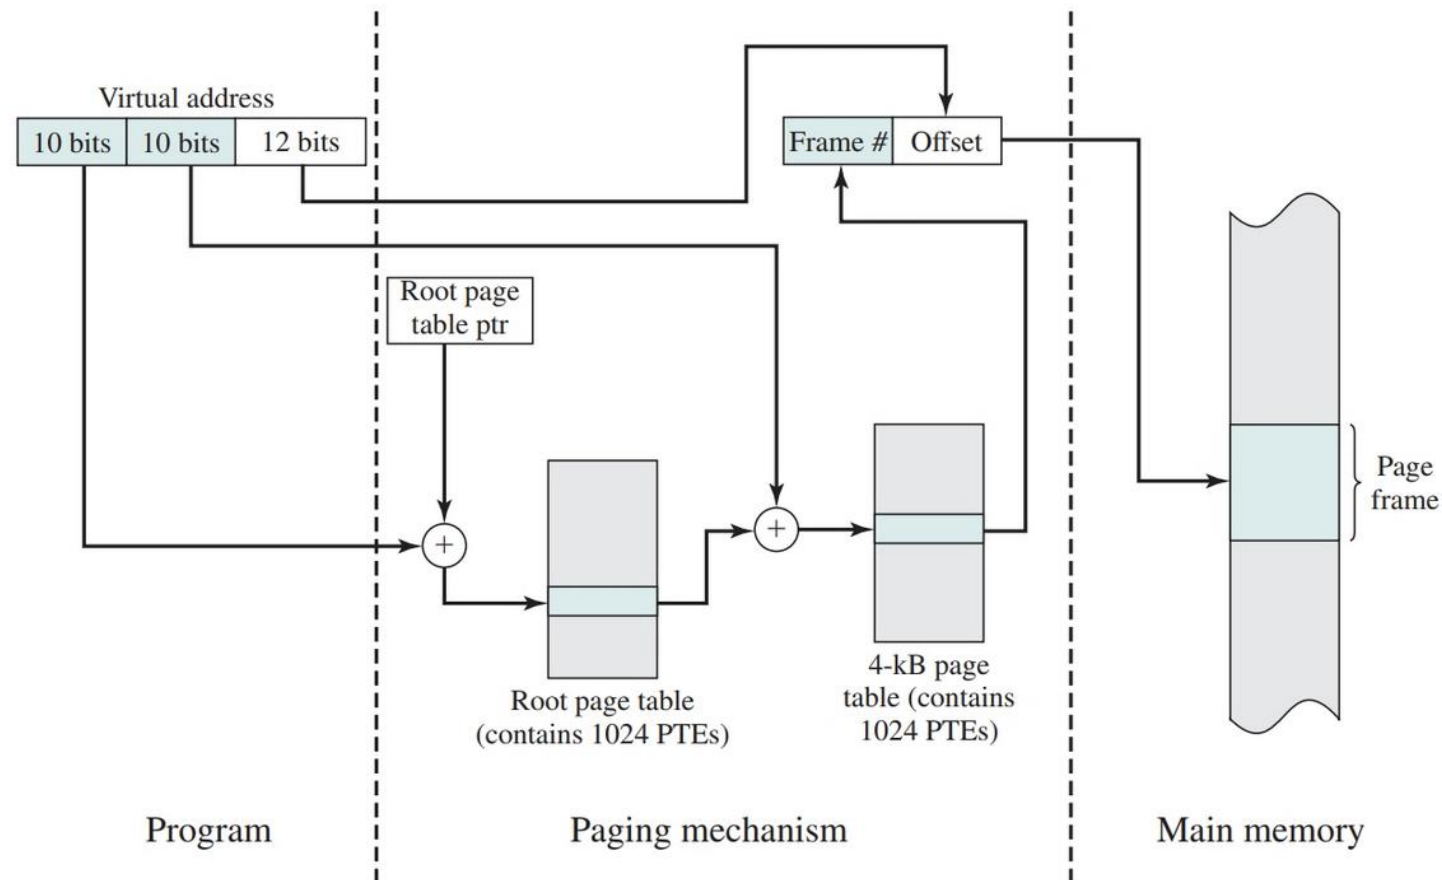


Figure 8.4 Address Translation in a Two-Level Paging System



ALTERNATIVE APPROACH: INVERTED PAGE TABLE

- Instead of having a page table entry of each possible virtual page, implement a hash table mapping the virtual pages in use to physical frames
 - Virtual page number is mapped to a hash value
 - The hash value is used to index the inverted page table
- Fixed proportion of real memory is needed for page tables regardless of the number of processes or virtual page numbers supported
 - There is one entry in the inverted page table for each real memory frame, rather than one per virtual page
 - More than one virtual page may map into the same hash value; A chaining technique used for managing the overflow



INVERTED PAGE TABLE STRUCTURE




Figure 8.5 Inverted Page Table Structure



INVERTED PAGE TABLE: EXAMPLE

- Virtual address ($n=8$):
0000 0111 0000 0001 1001




Figure 8.5 Inverted Page Table Structure



-
- Virtual address
 n bits
- Page # Offset
- n bits
- Hash function
- m bits
- Control bits
- Process
- Page # ID Chain
- | | | | | |
|---|----|---|--|-----------|
| | | | | 0 |
| | | | | |
| 4 | 12 | 6 | | i |
| | | | | |
| | | | | |
| 7 | 2 | - | | j |
| | | | | |
| | | | | |
| | | | | |
| | | | | $2^m - 1$ |
- Inverted page table
 (one entry for each
 physical memory frame)
- Frame # Offset
- m bits
- Real address

**HELSINGIN YLIOPISTO
HELSINGFORS UNIVERSITET
UNIVERSITY OF HELSINKI**



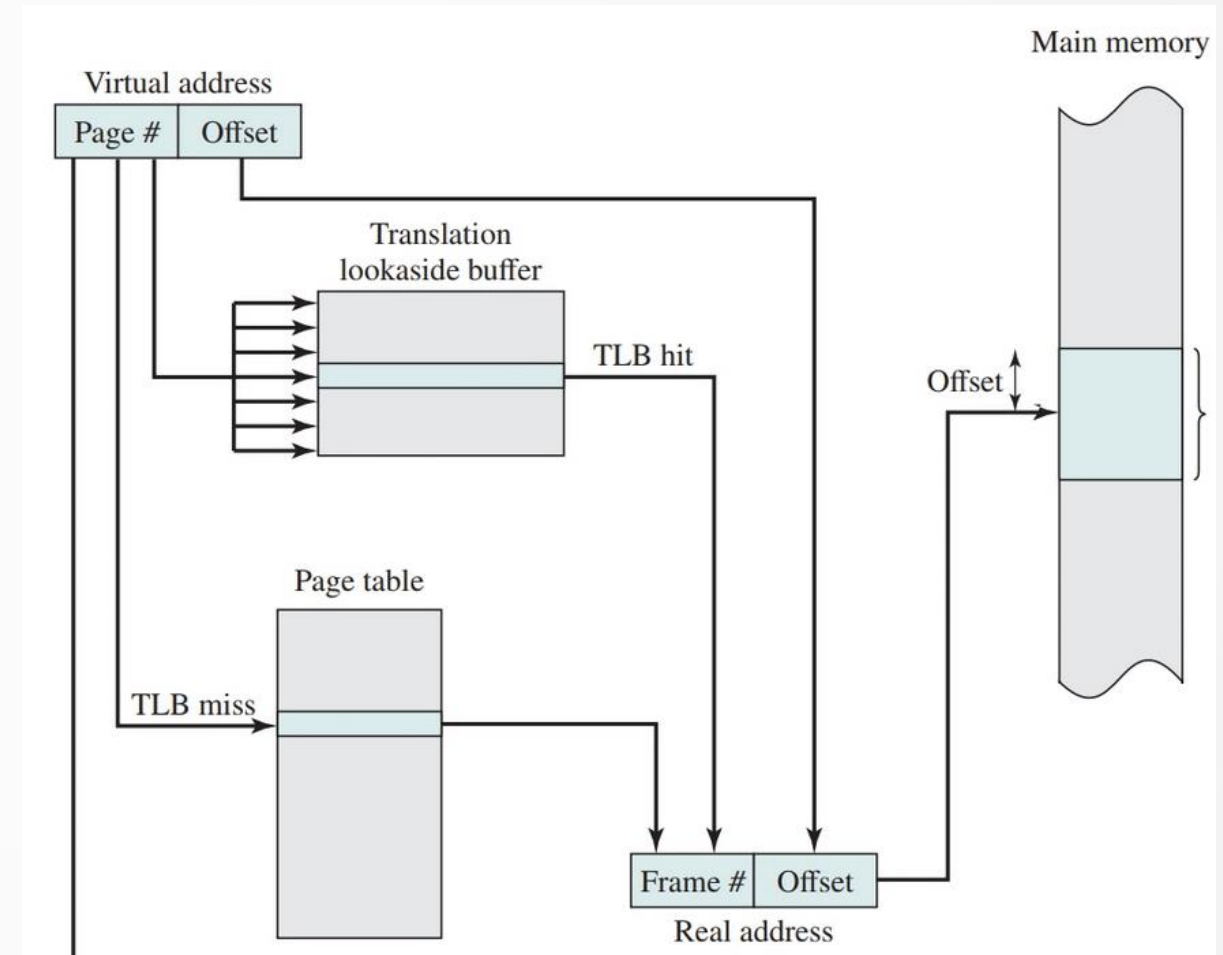
TRANSLATION LOOKASIDE BUFFER (TLB)

- Every virtual memory reference can cause **two physical memory accesses**
 - One to fetch the page table entry
 - Another to fetch the actual data
- To overcome the effect of doubling the memory access time, most systems employ a **special high-speed cache** called a **translation lookaside buffer (TLB)**
 - Hardware assistance!
 - Contains the page table entries that have been most recently used



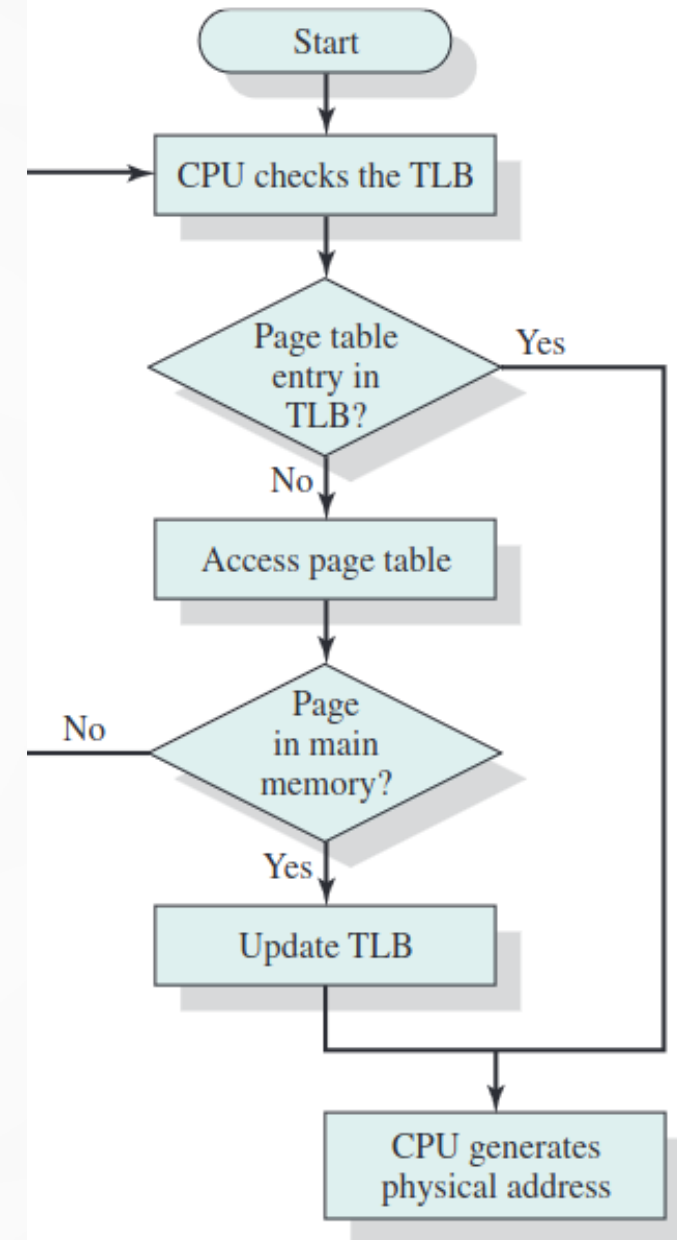
TRANSLATION LOOKASIDE BUFFER (TLB)

- TLB only contains some of the page table entries
 - We cannot simply index the TLB with page numbers
 - Each TLB entry must include the **page number** as well as the **complete page table entry**
- **Associative mapping:** Processor is equipped with hardware that allows to interrogate **simultaneously multiple TLB entries** to determine if there is a match on the page number





ADDRESS TRANSLATION WITH PAGING AND TLB





COMBINED SEGMENTATION AND PAGING

- In combined paging / segmentation system, the **user process address space is broken up into several segments**, and each segment is further broken up into several fixed-size pages (same size as frames in main memory)
- Segmentation is visible to the programmer
- Paging is transparent to the programmer
- From programmer's point of view: A virtual address consists of a segment number and segment offset
- From system's point of view: the segment offset is viewed as a page number and a page offset for a page specified within that segment



COMBINED SEGMENTATION AND PAGING

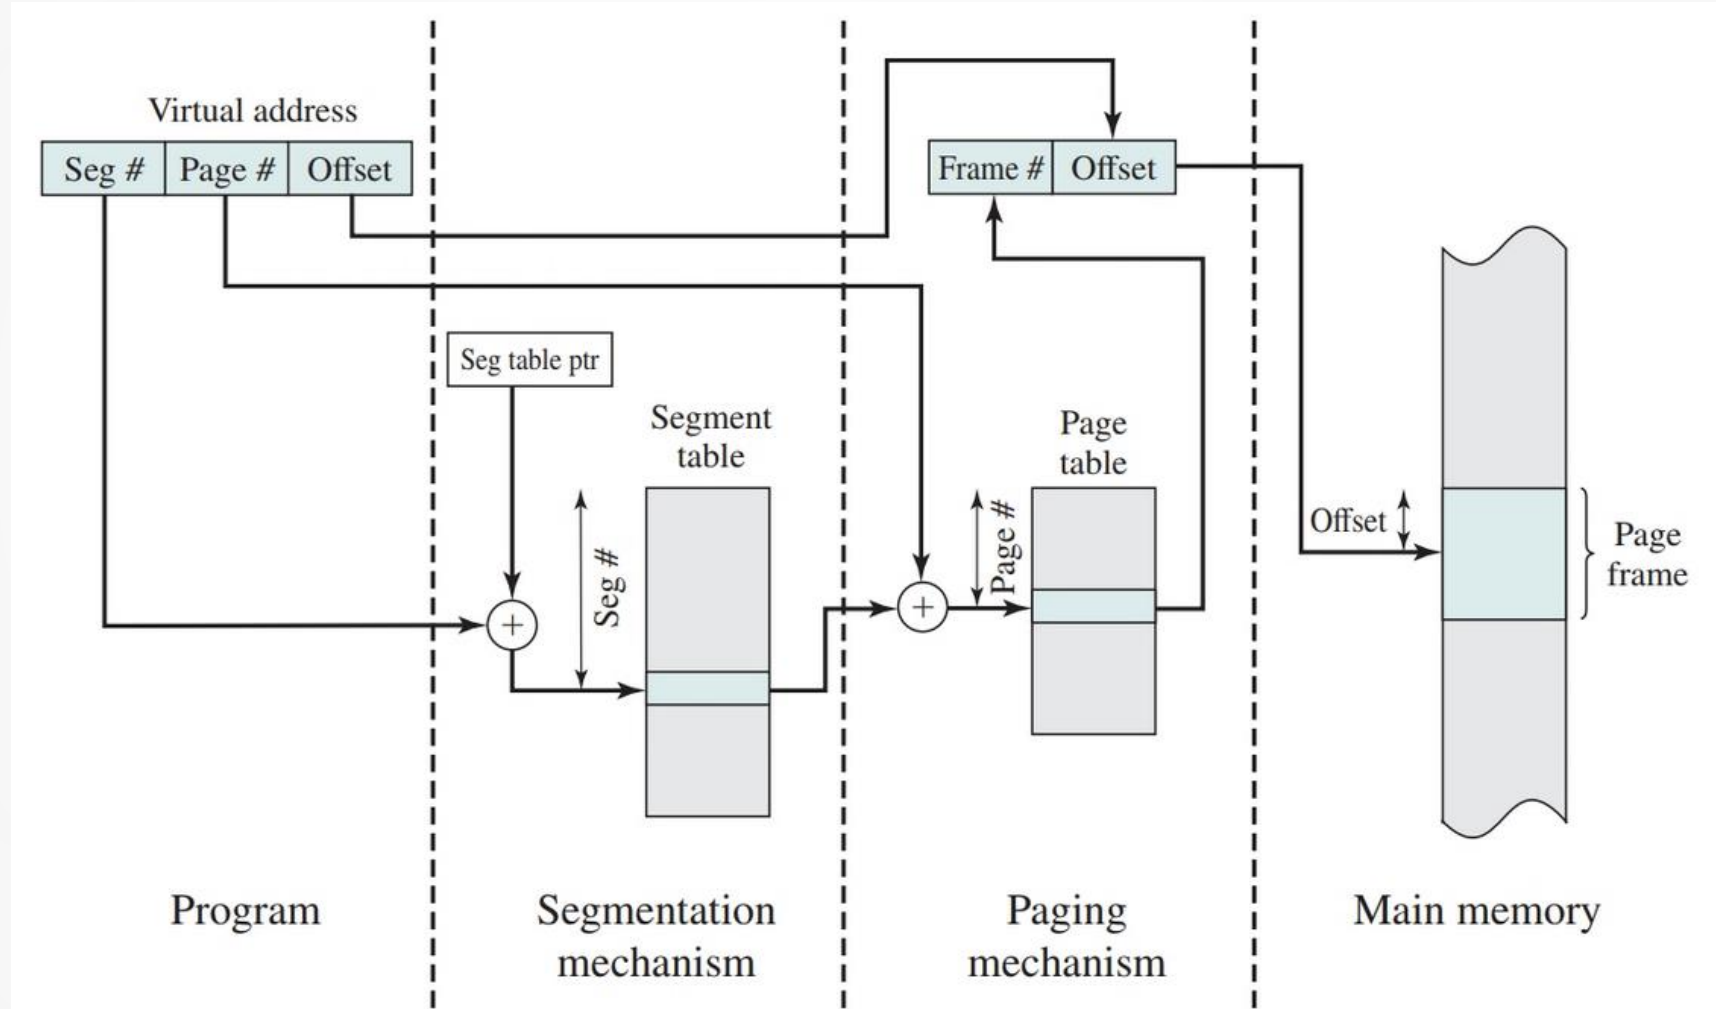


Figure 8.12 Address Translation in a Segmentation/Paging System



SEGMENTATION AND PAGING: EXAMPLE

- Virtual address:
1000 0110 1001 0010
 - Segment #: 10 = 2
 - Page #: 00 0110 = 6
 - Offset: 1001 0010
- 2nd entry in segment table gives the address of the page table
- 6th entry in the page table gives the frame number (assume this is 0001 0000 = 16)
- Physical address:
0001 0000 1001 0010

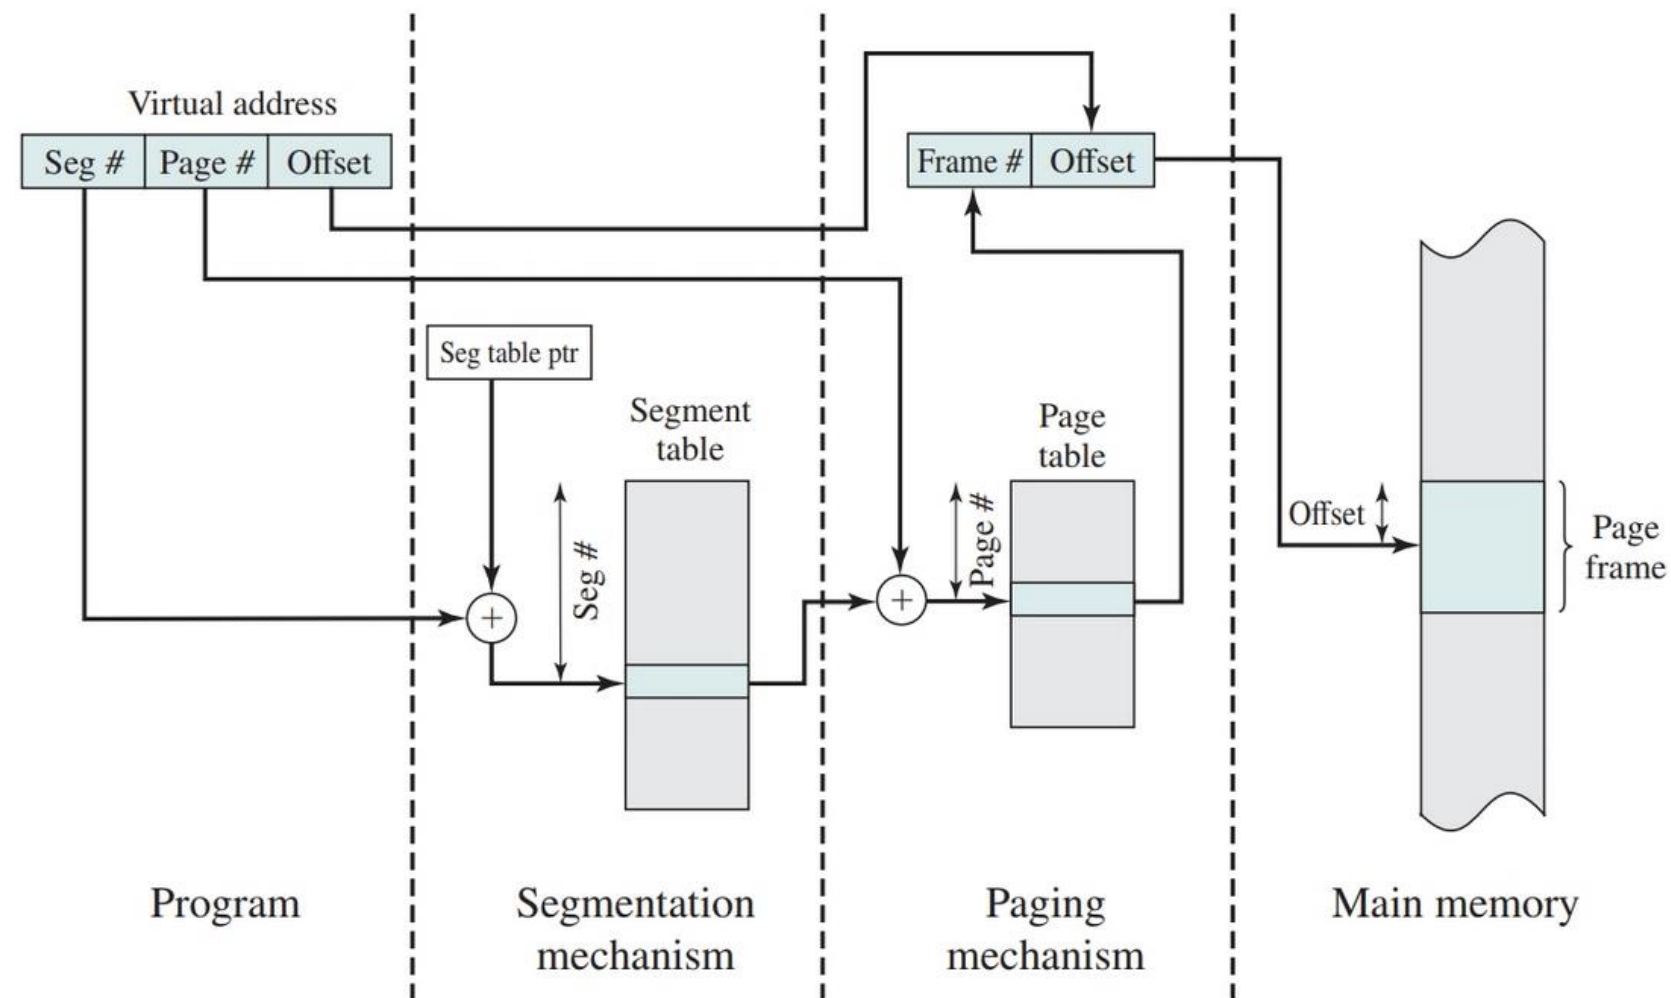


Figure 8.12 Address Translation in a Segmentation/Paging System



SUMMARY

- Memory management and virtual memory is one of the most important and complex tasks of the OS
- Memory is a resource that needs to be allocated to and shared among processes
- Memory management needs to allow **relocation** of processes, be **efficient**, and allow **protection and sharing** among processes
- Basic tools: Paging and segmentation (can also be combined!)
 - **Paging**: Small fixed-size pages
 - **Segmentation**: Pieces of variable size
- Implementation requires hardware assistance (e.g. TLB, address translation) and software support (e.g. segment tables, page tables)