



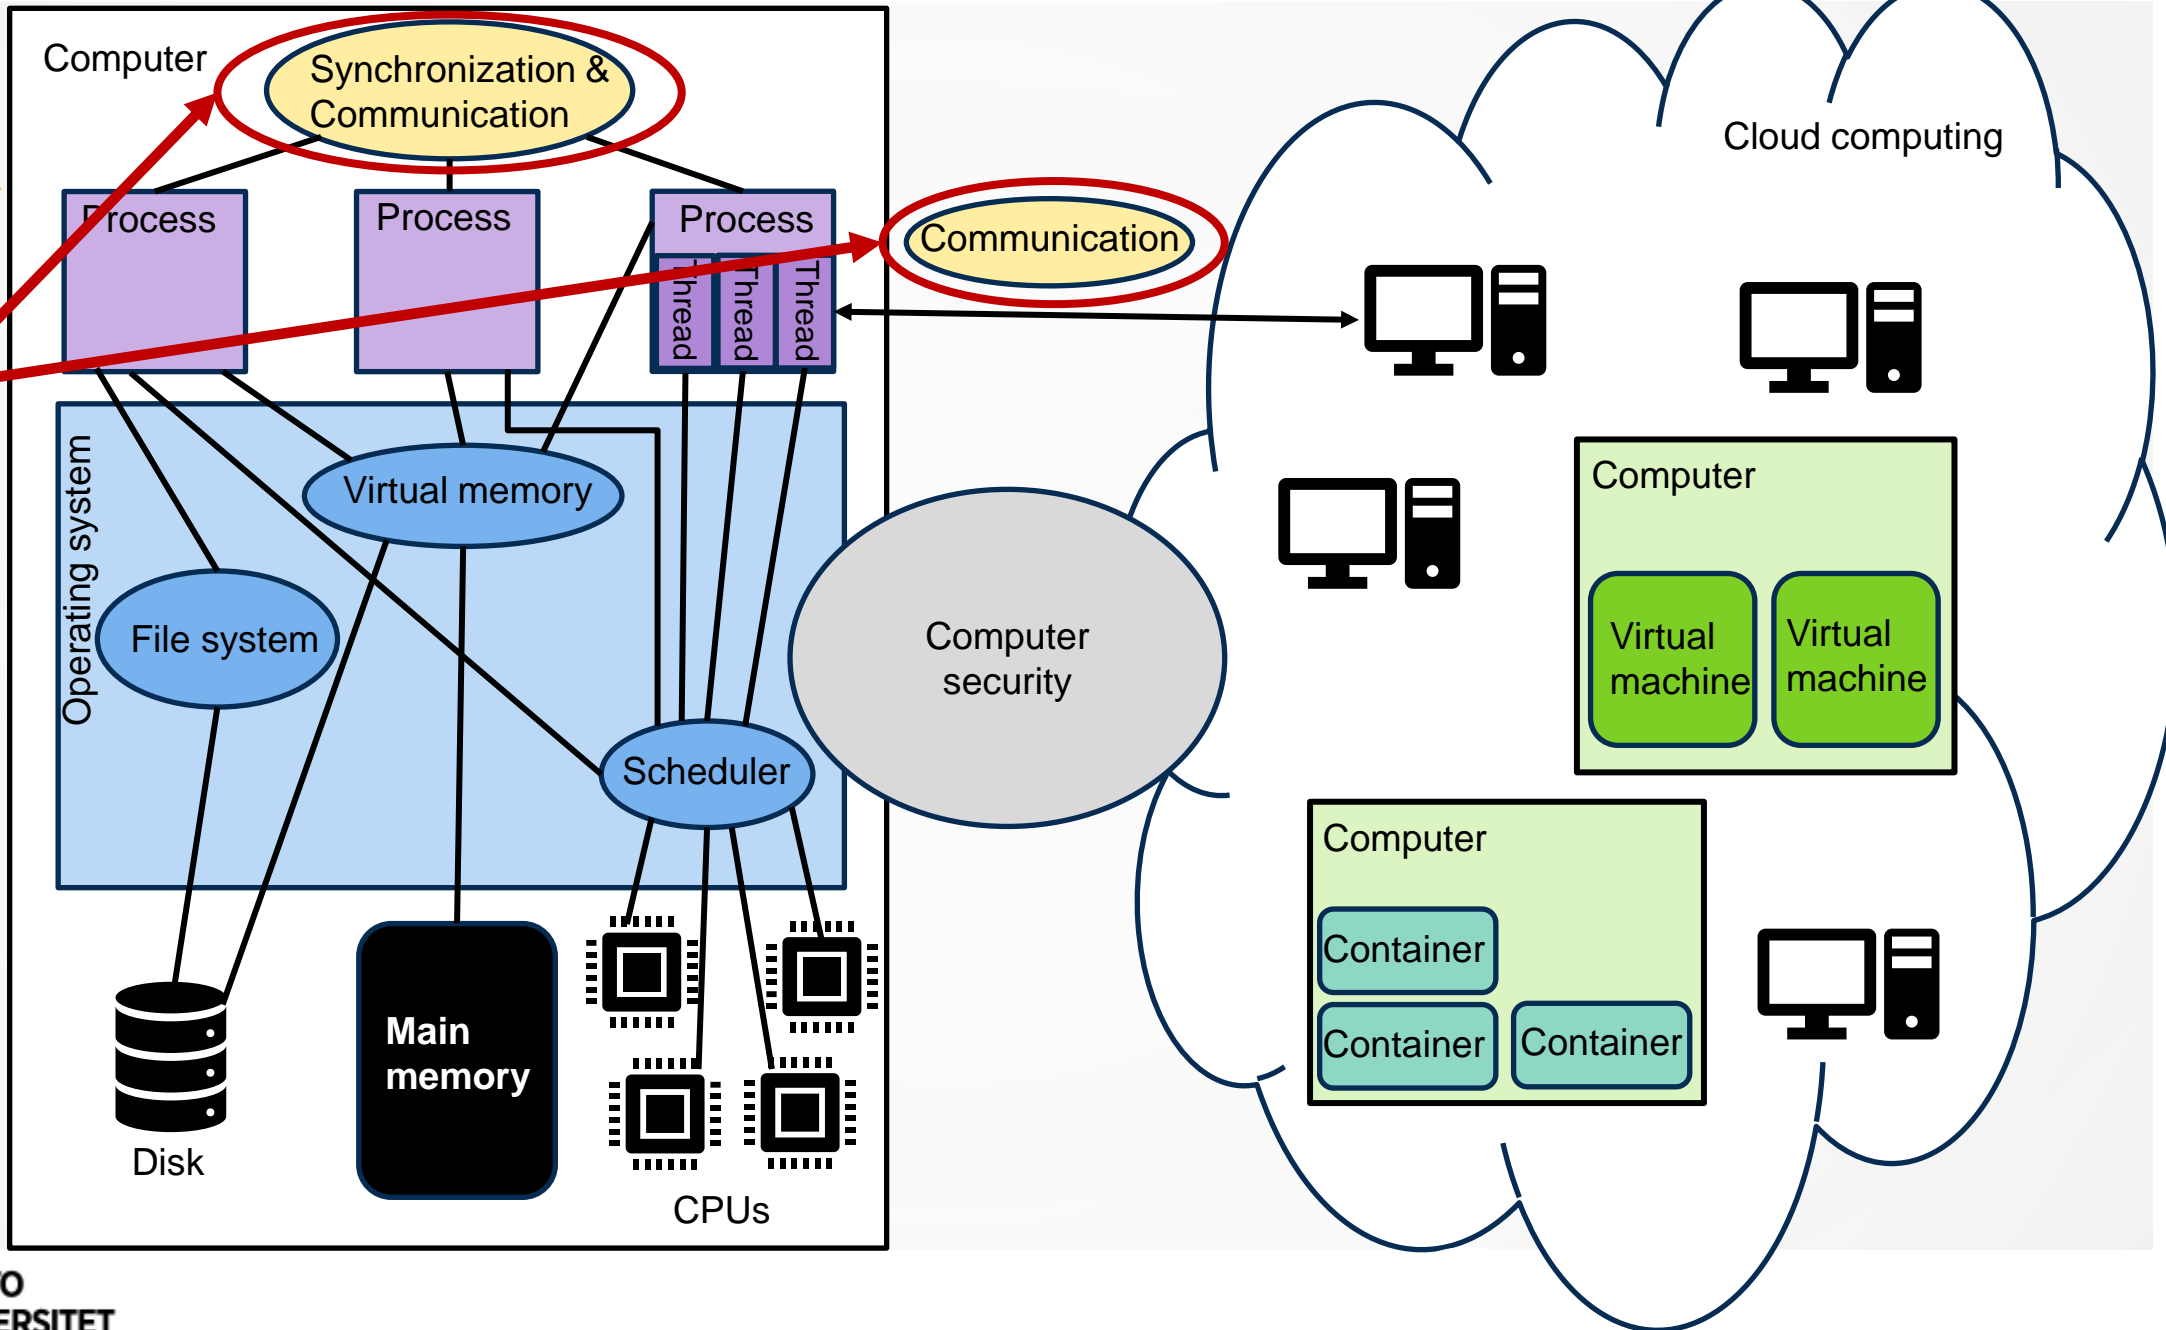
COMPUTING PLATFORMS

Interprocess communication:
Overview, background, alternatives, selection criteria

Mon 10.2.2025 Tiina Niklander



Today's
topic






CHECK YOUR COURSE EXAM REGISTRATION!

- Please remember to make separate registration to the course exam on time
- If you miss the registration deadline, you miss the exam



INTERPROCESS COMMUNICATION (IPC)

- Processes are isolated, but there is need for collaboration and coordination
- Several alternative methods and techniques that processes can use to exchange information
- Wikipedia: IPC are the mechanisms provided by an operating system for processes to manage shared data.



Coordinate actions
Synchronise operations
Exchange information

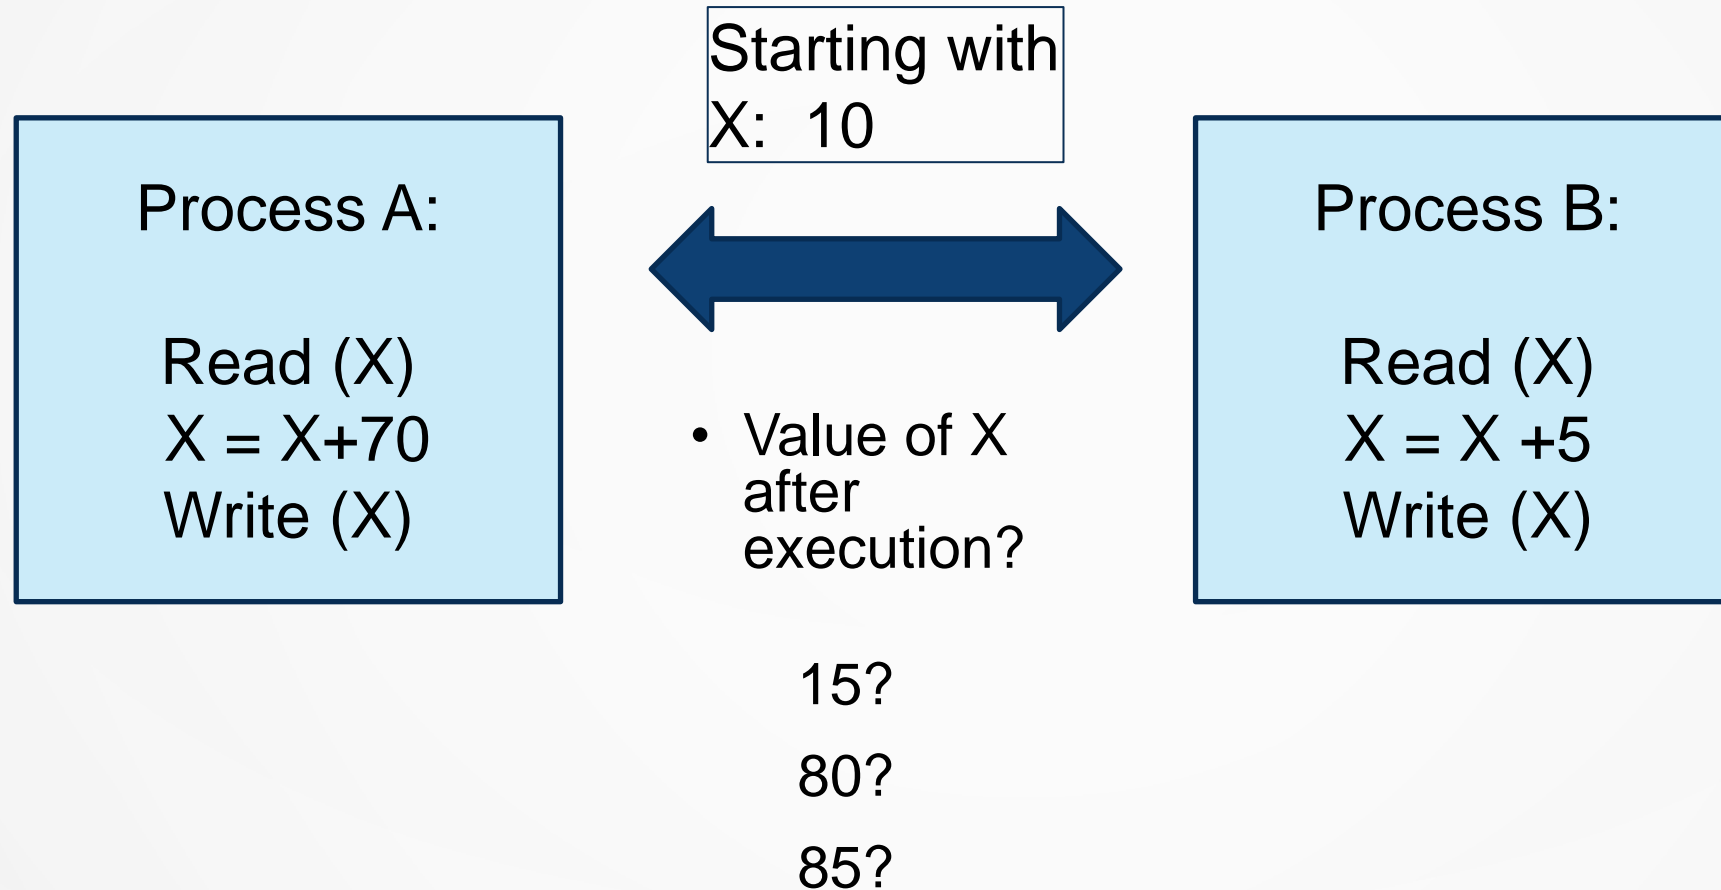


EXAMPLES

- Web-based services
 - Web browser needs to communicate with web server
 - Social media application needs to communicate with the server
- Parallel Computing
 - Used for example to model real-world events, like weather
 - One process can be running only on one core, but modern processors have multiple cores
- Others?



COORDINATION





TYPES OF COMMUNICATION

- Shared memory *Same computer*
 - Normal memory operations, but memory area shared with other process(es)
 - Read and write memory content
- Message-passing *Same or different computer*
 - Information between processes is passed using special messages
 - Send and receive messages
- Remote Procedure Call (RPC) / Remote Method Invocation (RMI) *Same or different computer*
 - One process can call a procedure or method of another process



TYPES OF COMMUNICATION, PART 2

- Signals
- Pipes / Named Pipes

Same computer

Must select one that is supported by API and OS

- Files

- Sockets
- Message queues

Same or different computer

Web services have their own mechanisms, that are built using sockets



SYNCHRONOUS AND ASYNCHRONOUS

Synchronous

- Both present at the same time
- Real-world analogy: telephone call
- Rendezvous

Asynchronous

- Can be present at different times
- Real-world analogy: text message, letter



BLOCKING VS NON-BLOCKING

- Communication is blocking when a process has to wait for the other process to answer before it can continue
- Synchronous communication is always blocking: The first one arriving to synchronisation point has to wait for the other to arrive
- Asynchronous communication can be blocking or nonblocking



SHARED MEMORY

- Fast – direct access to memory by both processes (on same computer)

- Critical issues:

- How to coordinate the operations by processes
- Exclusive access

➡ Next lectures

- Python examples:

https://docs.python.org/3/library/multiprocessing.shared_memory.html



SIGNALS

- Only in one computer
- Operating system specific
- Cannot be used to pass data, but just poke another process
- Needs process specific updates to signal handler on the receiving end
- Typically not used between user processes



FILES IN ONE COMPUTER

- Traditional way: One process writes to file, and then the other one reads from it
 - Very slow, if done frequently
- Memory-mapped files
 - Map the file to memory (like virtual memory page)
 - Operations on memory, but OS also updates the file to reflect the changes, but not immediately.
 - Can be used to implement shared memory
 - Python examples: <https://docs.python.org/3/library/mmap.html>



DISTRIBUTED FILE SYSTEMS

- Using same file in multiple computers, originally no concurrent access
- Distributed file systems (like NFS, SMB) allow file to be accessed from multiple computers
 - Network File System (NFS)
 - Developed by Sun microsystems in 1984
 - Used still in Unix-based systems like Linux
 - Server Message Block (SMB)
 - Developed at IBM in 1983
 - Used still in Windows to share files and other resources (like printers)
- Cloud-based services



PIPES AND NAMED PIPES (A.K.A. FIFO)

- Typical use on commandline: forward output of one process as input to another
 - `ls | head 10`
- Always unidirectional – One process writes and the other one reads
 - Stdin and stdout are examples of pipes used in linux
- Using pipes in python: <https://docs.python.org/3/library/subprocess.html> or <https://docs.python.org/3/library/multiprocessing.html>
- Named pipes (fifos): <https://docs.python.org/3/library/os.html>
 - Named pipe is created with command `os.mkfifo`
 - Used a lot like files



MESSAGE PASSING

Message-passing has multiple different alternatives

- Sockets are typically used with separate computers, but can be used in one computer
- Message queues have different implementations for use in one or more computers

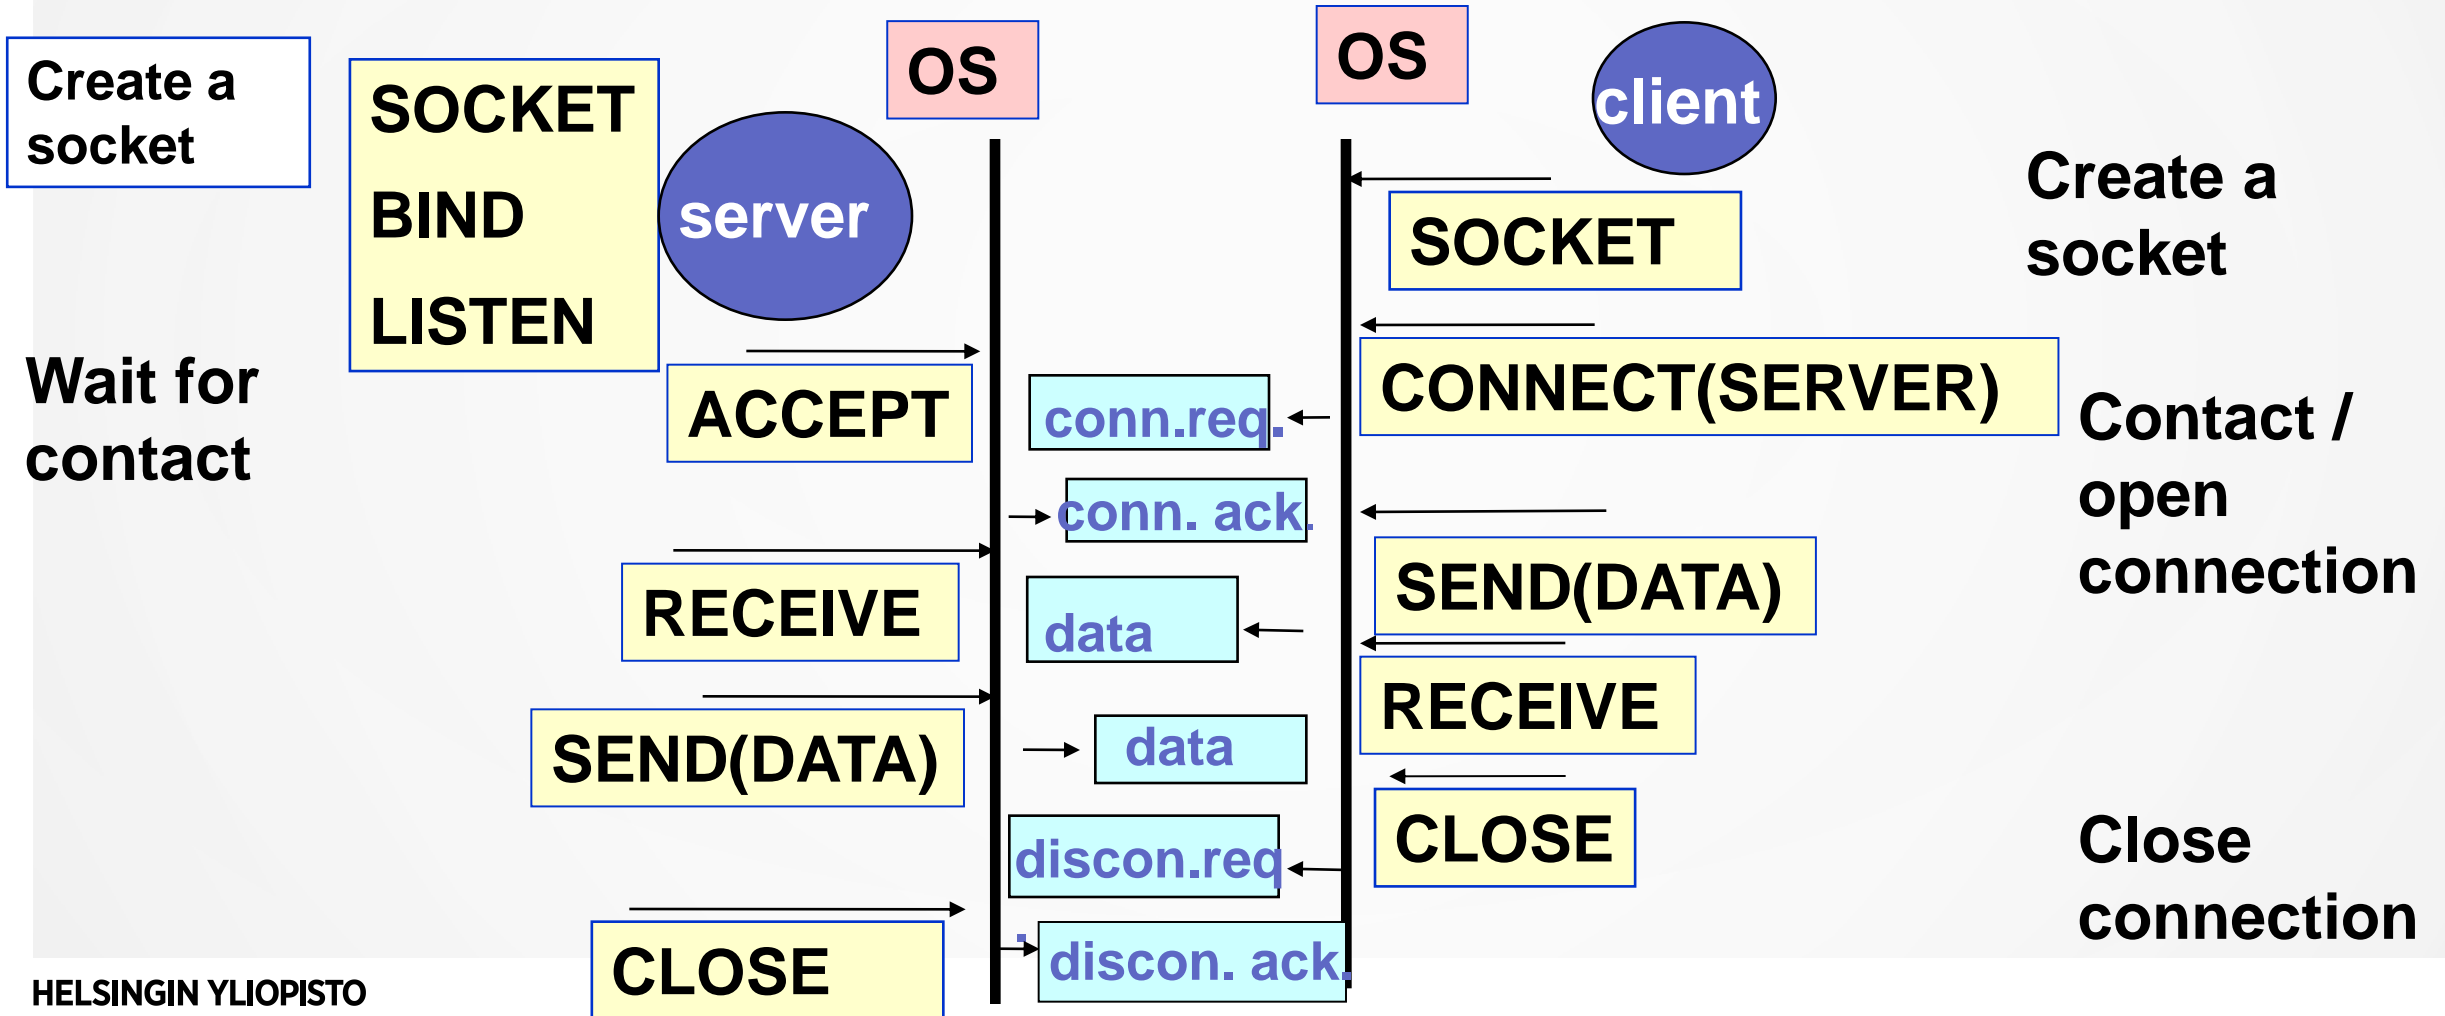


SOCKETS

- Sockets are the basic communication interface in Internet-based applications
 - Can be used in one computer also (IP 127.0.0.1)
- Note: Underlying all modern Internet-related communication mechanisms, even if not used as such
- Require both parties to be active at the same time (synchronous), but handling one message can be synchronous or asynchronous
- Using in python: <https://docs.python.org/3/library/socket.html> with or without <https://docs.python.org/3/library/ssl.html>
- Covered in Network Programming -course




TCP SOCKETS (AND MESSAGES)





MESSAGE QUEUES

- Message queues provide asynchronous mechanism to pass information from one process to another
 - One process writes / sends / produces a message to the queue
 - Other process reads / receives / consumes a message from the queue
 - Access to the queue needs some synchronization (like locks)  Next lectures
- One computer using python: <https://docs.python.org/3/library/multiprocessing.html>
- Multiple computers: must have a queue manager (process) or use a message queue middleware like RabbitMQ, Apache Kafka, ...



MESSAGE QUEUES

- Stallings calls message queues (MQ) mailboxes
 - unidirectional
- Process attach itself to MQ, where it gets/reads/consumes messages
- Multiple processes consuming messages in one mail box
 - One message to all of them or
 - Each message to just one process
 - first reader,
 - select based on message type, ...




Figure 5.21 Indirect Process Communication

Stallings



REMOTE PROCEDURE CALL (RPC)/ REMOTE METHOD INVOCATION (RMI)

- Used to execute one procedure or method on a remote computer
- The caller process making the call is client
- The callee process executing the procedure is server

- Client needs to have a stub that
 - Can pack the argument to a message
 - Send the message to server
 - Wait for the reply
 - Unpack the result from message

- Server needs to have a stub that
 - Can wait for the incoming message
 - Unpack the arguments
 - Make the actual procedure call
 - Pack the result to message and send it



RPC

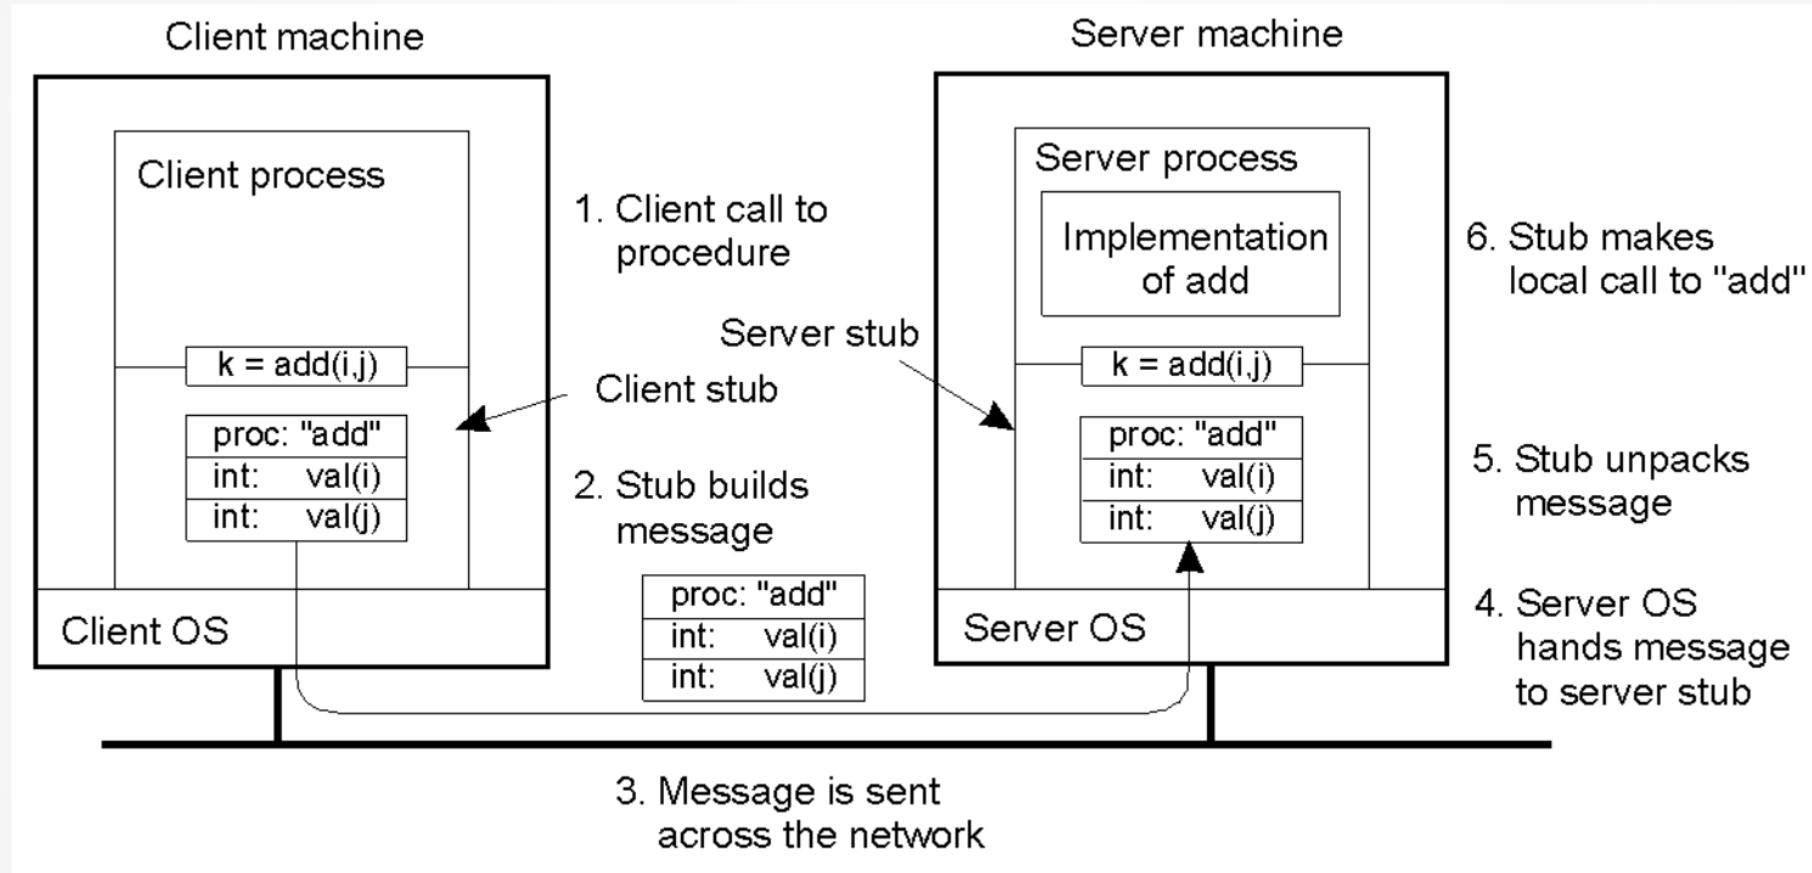


Figure from Tanenbaum and van Steen: Distributed Systems



DISCUSSION: FIRST PAIRS, THEN LARGER GROUPS

- Features

Synchronous
Asynchronous

Blocking
Non-blocking

and methods:

Shared-memory
Signals
Pipes
Files
Sockets
Message queues

- Explain the listed features and methods
- Which features describe which methods? Why / why not?
- If you have already used some of these, explain when and how
- If you have not used, think about situations when you might use each of them



MESSAGE PASSING VS SHARED MEMORY

- Shared memory
 - Used in one computer, but
 - Distributed Shared memory solutions have been studied

https://en.wikipedia.org/wiki/Distributed_shared_memory

- Message passing
 - Must identify the recipient
 - Must define protocol including message formats
 - Value encodings and marshalling



MESSAGE PASSING: ADDRESSING

- Direct addressing
 - Sockets: (*IP address, port number*)
 - IP address identifies the destination computer (and source computer)
 - Local computer can always use IPv4 address 127.0.0.1 (loopback address)
 - Other computer's IP address must be known or searched (e.g. DNS, directory service, etc)
 - Port number identifies the socket (and receiving/sending process)
- Indirect addressing
 - Message queues are identified instead of processes
 - Messages are sent to MQ, not to the process
 - Access to MQ can be exclusive (just one process) or shared (multiple processes)



MIDDLEWARE (VÄLIOHJELMISTO)

- Wikipedia: Middleware is a type of computer software program that provides services to software applications beyond those available from the operating system.
- IETF (in 2000): services found above the transport (i.e. over TCP/IP) layer set of services but below the application environment
- One example: Object Request Brokers (ORB)
- Message Oriented Middleware (MOM)
 - Create uniform interface for application
 - Use message brokers that
 - store (buffer), route, and/or transform messages

See Wikipedia for more information:

- Message brokers
- Message Oriented Middleware



SECURITY ISSUES

- Who can access message content?
 - Only the sending and receiving processes?
 - Also message oriented middleware?
 - Also other message passing entities between sender and receiver?
- Use encryption and proper key management
 - Message protected, but
 - All marshalling between different mappings have to be done by sender and receiver
 - Data types and sizes, Big-Endian, Little-Endian, different character encoding, ...
- What is the priority?



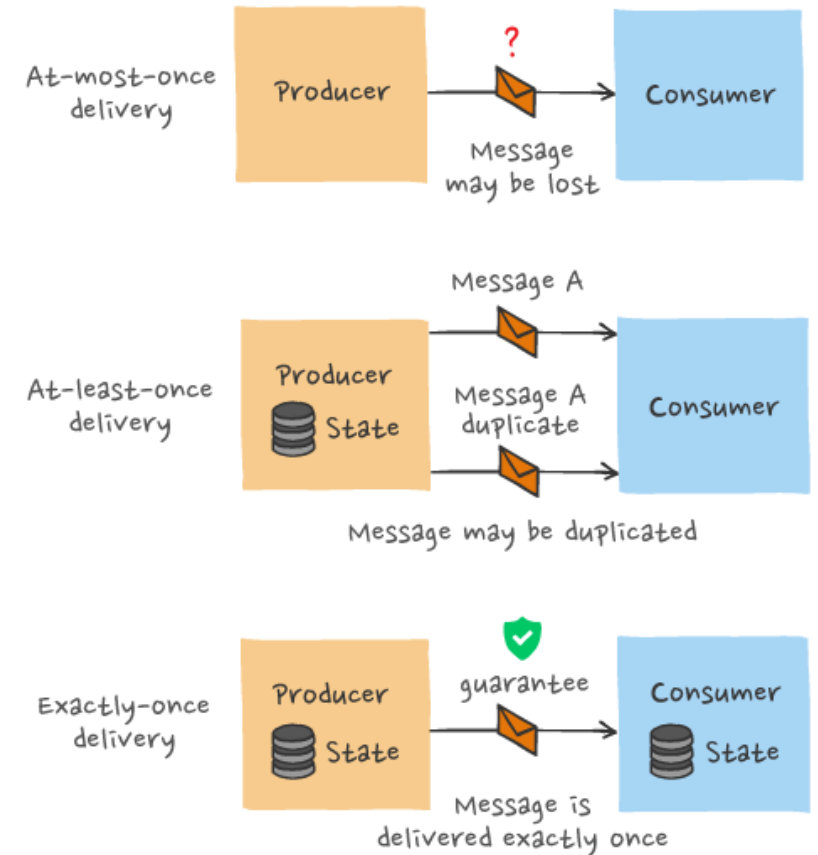
MESSAGE PASSING AND COMMUNICATION SEMANTICS

- Typical assumption: All messages transferred in the sending order
- However, a message in transit can
 - corrupt
 - be lost
 - duplicate
 - change order with other messages
- Typical solution: Detect and Retransmit
 - How many times the message is received?



MESSAGE DELIVERY SEMANTICS

- At-most-once semantics
 - A sent message might not be delivered
 - No message is delivered multiple times
- At-least-once semantics
 - A sent message is (eventually) delivered
 - A message might be delivered multiple times (duplicates allowed)
- Exactly-once semantics
 - Ideal target, but not always possible to reach





HOW TO SELECT?

- Based on communication style:
 - One-to-one, many-to-one, one-to-many, many-to-many
 - Synchronous vs asynchronous, blocking vs non-blocking, direct vs indirect
- Based on availability of methods
 - Not all OSs or programming languages provide all alternatives
- Based on number of computers
 - Single computer, multiple networked computers
- Based on efficiency / performance / security / semantics goals
- Based on ...



SUITABLE READING

- Inter-process communication: https://en.wikipedia.org/wiki/Inter-process_communication
- Other pages linked from that
 - Message passing
 - Shared memory
 - Pipe, socket, message queue
- Stallings has just few pages and use slightly different terminology



TRANSIENT AND PERSISTENT MESSAGE PASSING

Transient

- Message in system only as long as both processes are executing
- Pipes and sockets are transient

Persistent

- Sent message is stored in the system as long as it takes to deliver it to the recipient
- Files and message queues are persistent



EXTRA DISCUSSION TOPICS

- Time overlap?
 - Processes must be running same time or can be running different times
- Speed of the mechanism
 - Which might be slowest? Why?
- How much data can be passed conveniently in one contact?
 - Differences? Can you order mechanisms based on the data amount?

Shared-memory
Signals
Pipes
Files
Sockets
Message queues