

Computing platforms (Spring 2025)

week 3

Exercise 1

Address Translation: Segmentation

32-bit address space, leftmost 4 bits represent segment number and remaining 28 bits offset within that segment. The segmentation table for a process is as follows:

Index	Base Address	Segment Length
0	0x00001000	0x0500
1	0x00020000	0x5000
2	0x00030000	0x10000
3	0x00045000	0x5000

(a) Translating Virtual Address 0x20001111

The virtual address can be split into:

Segment Number	Offset
0x2 (Segment 2)	0x0001111

For segment 2:

Segment	Base Address	Segment Length
2	0x00030000	0x10000

Since 0x0001111 is within the segment length 0x10000, the physical address is:

$$\text{Physical Address} = \text{Base Address} + \text{Offset} = 0x00030000 + 0x0001111 = 0x00031111$$

Answer for (a): The virtual address 0x20001111 translates to the physical address 0x00031111.

(b) Translating Virtual Address 0x00001000

The virtual address can be split into:

Segment Number	Offset
0x0 (Segment 0)	0x0001000

For segment 0:

Segment	Base Address	Segment Length
0	0x00001000	0x0500

Since the offset 0x1000 exceeds the segment length 0x0500, segmentation fault occurs.

Answer for (b): The virtual address 0x00001000 is invalid because the offset exceeds the segment's size, causing segmentation fault.

Exercise 2

Address Translation: Paging

We have a system with:

- 32-bit virtual addresses
- 4 KB pages ($4 \text{ KB} = 2^{12}$ bytes)
- A single-level page table

(a) Layout of Virtual Addresses

Since each page is 2^{12} bytes in size, the **offset** occupies the lower 12 bits of the address. The remaining upper 20 bits specify the **page number**. Thus, a 32-bit virtual address is divided as follows:

$$\begin{array}{c} \underbrace{\text{Page Number}}_{20 \text{ bits}} \mid \underbrace{\text{Offset}}_{12 \text{ bits}} \\ \text{[5 hex digits] [3 hex digits]} \end{array}$$

(b) Translating Virtual Address 0x00001234

Assume a process has single-level page table shown below:

Index	Frame Number
0	3
1	2
2	7
3	1

Step 1: Split the Virtual Address.

$$0x00001234 \longrightarrow \begin{cases} \text{Page Number} = 1 \\ \text{Offset} = 0x234 \end{cases}$$

Step 2: Look Up the Frame. Using page number = 1 \longrightarrow Frame Number = 2

Step 3: Construct the Physical Address. Each frame is also $4 \text{ KB} = 2^{12}$ bytes hence frame number 2 starts at

$$\underbrace{2 \times 2^{12}}_{\text{Frame Number} \times \text{Frame Size}} = 2 \times 0x1000 = 0x2000.$$

and add offset 0x234:

$$\text{Physical Address} = 0x00002000 + 0x00000234 = 0x00002234.$$

Exercise 3

Virtual Memory with Paging and TLB

We have:

- 32-bit virtual addresses
- 4 KB pages ($4 \text{ KB} = 2^{12} \text{ bytes}$)
- A single-level page table and a Translation Lookaside Buffer (TLB)

Page Table and TLB Contents

single-level page table:

Page table				TLB				
Index	Frame Number	M	P	Index	Page #	Frame #	M	P
0	3	0	1	0	3	1	1	1
1	2	0	1	1	0	3	0	1
2	7	0	0	2				0
3	1	1	1	3				0

- M = 1 indicates the cached page is modified.
- P = 1 indicates the TLB entry is valid (present).
- Entries 2 and 3 are either empty or invalid (P = 0).

Address Format (32-bit, 4 KB Pages)

A 32-bit address is split into:

$$\underbrace{\text{Page Number}}_{20 \text{ bits}} \mid \underbrace{\text{Offset}}_{12 \text{ bits}}$$

Hence, the lower 12 bits (3 hex digits) are the *offset*, and the upper 20 bits are the *page number*.

(a) Translating Address 0x00001123

Step 1: Split the Address.

0x00001123

- Offset = 0x123 (the lowest 12 bits)
- Page Number = 0x1 (the upper 20 bits)

Step 2: TLB Lookup. We check whether the TLB has an entry for Page 1.

- Index 0 has Page 3.
- Index 1 has Page 0.
- Index 2,3 are invalid.

No match \Rightarrow **TLB miss.**

Step 3: Page Table Lookup. From the Page Table:

Page 1 \longrightarrow Frame 2, $M = 0$, $P = 1$

Since $P=1$, the page is in memory (no page fault).

Step 4: Construct Physical Address. Frame 2 means the base of this frame is:

$$2 \times 2^{12} = 0x2000.$$

Adding the offset 0x123:

$$\text{Physical Address} = 0x2000 + 0x0123 = 0x2123.$$

Step 5: Update TLB. We add a new TLB entry for Page 1:

Page No. = 1,
Frame No. = 2,
 $M = 0$ (no writes so far),
 $P = 1$ (valid entry).

This typically occupies an empty TLB slot (say Index 2).

The virtual address 0x00001123 translates to 0x00002123. The TLB is updated to include Page 1 \rightarrow Frame 2.

(b) Translating Address 0x00002333

Step 1: Split the Address.

0x00002333

- Offset = 0x333 (lowest 12 bits)
- Page Number = 0x2 (upper 20 bits)

Step 2: TLB Lookup. We check for Page 2 in the TLB:

- Index 0 has Page 3,
- Index 1 has Page 0,
- Index 2 (updated above) has Page 1,
- Index 3 is invalid/empty.

No match \Rightarrow **we got TLB miss again.**

Step 3: Page Table Lookup. From the Page Table:

Page 2 \longrightarrow Frame 7, $M = 0$, $P = 0$

$P=0$ means the page is not present in main memory \Rightarrow **page fault.**

Step 4: Handle Page Fault. The OS must load Page 2 from disk into a free frame. Assume the OS chooses some free frame F . After loading:

Page 2 now mapped to Frame F ,
 $M = 0$ (still no writes), $P = 1$ (now present).

The page table is updated accordingly:

Page Table[2] \leftarrow (Frame F , $M = 0$, $P = 1$).

Step 5: Construct Physical Address. Once the page is loaded, the physical base address is:

$$F \times 2^{12},$$

and we add the offset 0x333 to get the final physical address.

Step 6: Update TLB. We add a valid entry for Page 2 in the TLB:

Page 2 \longrightarrow Frame F , $M = 0$, $P = 1$.

The virtual address 0x00002333 causes a page fault. After the OS loads Page 2 into some free frame F , the physical address is $(F * 0x1000) + 0x333$. The page table and TLB entries for Page 2 are updated to reflect Frame = F and $P=1$.

Exercise 4

Page Replacement Comparison

We have:

- Main memory of 3 frames
- Reference stream: 7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 2
- We count page faults only after the first 3 references have initialized the frames.

Ref.	7	0	1	2	0	3	0	4	2	3	0	3	2
FIFO	7	7	7	2	2	2	2	4	4	4	0	0	0
	-	0	0	0	0	3	3	3	2	2	2	2	2
	-	-	1	1	1	1	0	0	0	3	3	3	3
				F		F	F	F	F	F	F		
LRU	7	7	7	2	2	2	2	4	4	4	0	0	0
	-	0	0	0	0	0	0	0	0	3	3	3	3
	-	-	1	1	1	3	3	3	2	2	2	2	2
				F		F		F	F	F	F		
Clock	7	7	7	2	2	2	2	4	4	4	4	4	4
	-	0	0	0	0	0	0	0	2	2	2	2	2
	-	-	1	1	1	3	3	3	3	3	0	3	3
				F		F		F	F		F		
OPT	7	7	7	7	7	7	7	2	2	2	2	2	2
	-	0	0	0	0	3	3	3	3	3	3	3	3
	-	-	1	2	2	2	0	0	0	0	4	4	4
				F		F		F			F		

Table 1: Comparison of FIFO, LRU, OPT, and Clock. “F” marks a fault.

Fault Counts and Miss Rates

- **FIFO:** 7 faults over 10 references \Rightarrow 70% miss rate.
- **LRU:** 6 faults \Rightarrow 60% miss rate.
- **Clock:** 4 or 5, depending on the initial hand and use-bit detail.
- **Optimal:** 5 faults if you consider extended reference stream $(\dots, 1, 2, 0, 1, 7, 0, 1)$ to decide which page is used farthest in the future.

Exercise 5

Consider “matrix-rowwise.py” and “matrix-columnwise.py”

(a) Memory Organisation of the Two-Dimensional Table

In Python “list of lists” is stored as a top-level list, each element of which is a separate Python list object. Therefore, only the references in row j are contiguous within that row’s list structure; there is no guarantee that row $j+1$ is adjacent to row j in memory. Each integer is also a Python object on the heap, so even within a single row, the values themselves need not be stored contiguously as raw bytes.

(b) Run both of the codes and measure how much time their execution takes

```
$ time python3 ./matrix-columnwise.py
171475500000
```

```
real    0m6.811s
user    0m6.607s
sys     0m0.199s
```

vs.

```
$ time python3 ./matrix-rowwise.py
171475500000
```

```
real    0m3.967s
user    0m3.778s
sys     0m0.185s
```

Columnwise reading in this case is slower at reading the memory.

(c) Explain the difference in the running times.

Row-wise traversal (`array[j][i]`) accesses elements from same inner list meaning it follows consecutive memory addresses for that row’s pointers, improving cache locality. **Column-wise traversal** (`array[j][i]`) accesses elements from different lists (one list per row), meaning each lookup requires fetching a new row’s list reference, which may be far apart in memory. This leads to more cache misses.

Comparison to NumPy Arrays If the same operation were done using a **NumPy array**, where elements are stored in a **contiguous block of memory**, cache strides would be effective, and column-wise traversal would be much faster. This demonstrates why NumPy is preferred for numerical computing in Python.

```
import numpy as np

cols = 7000
rows = 7000
array = np.arange(cols, dtype=np.int64).repeat(rows).reshape(rows, cols)

# axis 0 or 1 control column/row reading
summa = sum(np.sum(array, axis=0))

print(summa)
---
```

```
$ time python3 ./numpy-arr.py
171475500000

real    0m0.129s
user    0m0.075s
sys     0m0.049s
```